



Tin hoc Nha Uyên

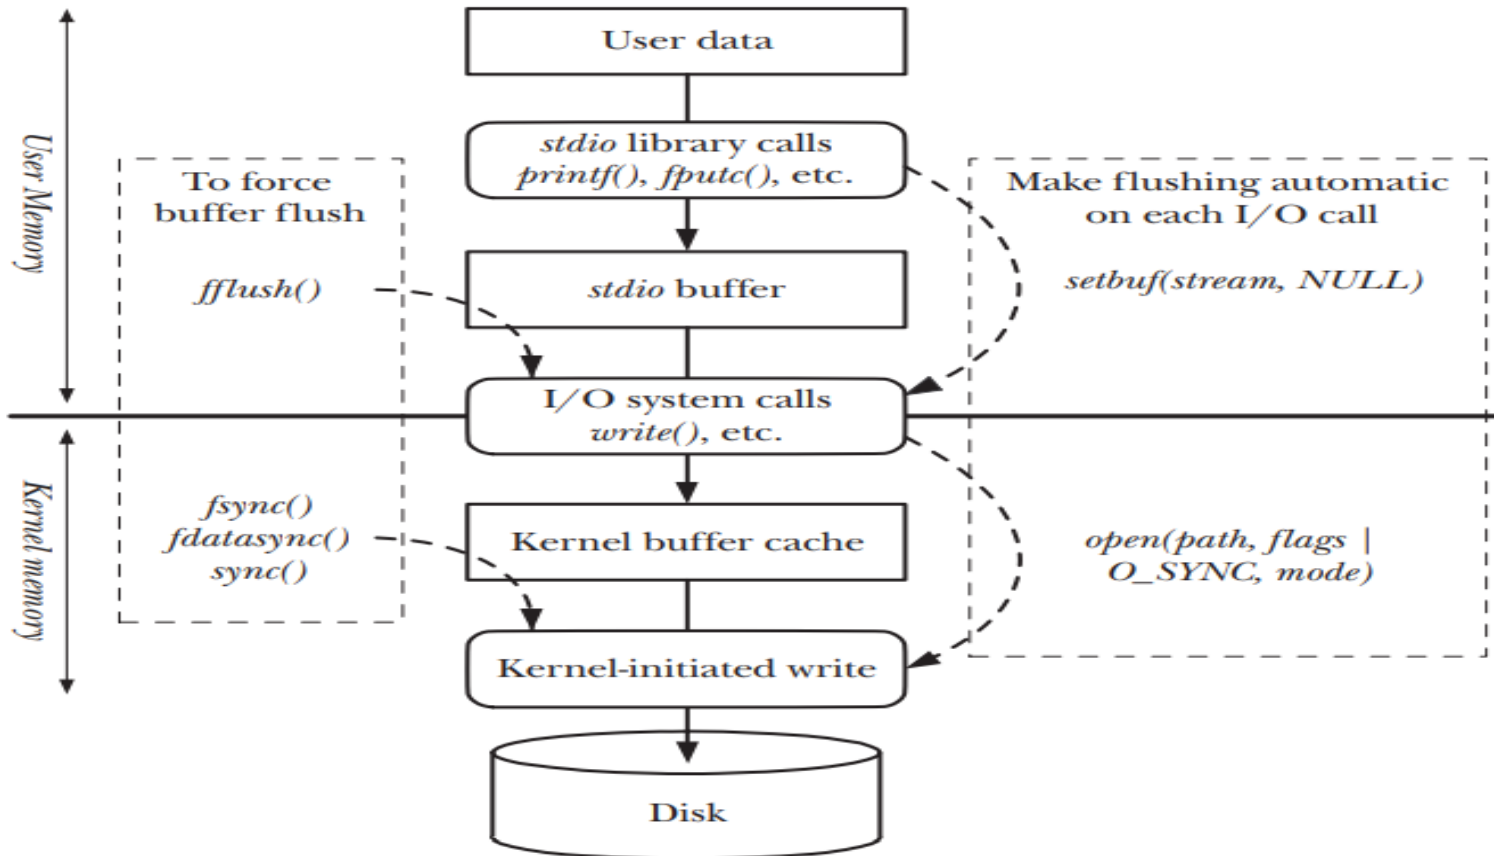
File in Linux

TIN HOC NHÃ UYÊN 10/2018

Content

- ▶ Overview File in Linux
- ▶ File I/O basic method
- ▶ File I/O atomic and race condition
- ▶ File control operation
- ▶ Relationship between File Descriptor and Open File
- ▶ File I/O some advance method
- ▶ File buffering and stdio library buffering

Overview File in Linux



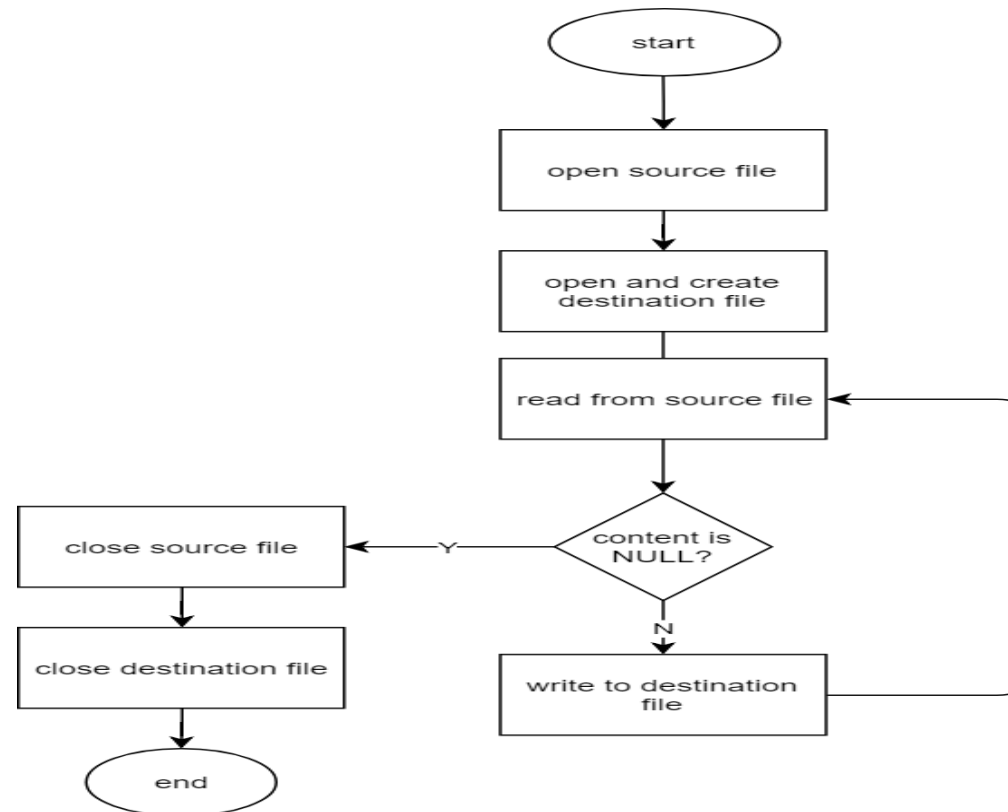
File I/O - Overview

- ▶ File is central to the UNIX philosophy
- ▶ All system calls for performing I/O referring to a file descriptor
 - ▶ Is nonnegative integer, start from 0
 - ▶ File descriptor are used to refer to all types of file, pipe, FIFOs, socket, terminals, devices
 - ▶ Each process has it own set of file descriptor
- ▶ Most of program expect to able to use three standard file descriptor
 - ▶ Running process inherit copies of the shell opened file descriptors

File descriptor	Purpose	POSIX name	stdio stream
0	standard input	STDIN_FILENO	<i>stdin</i>
1	standard output	STDOUT_FILENO	<i>stdout</i>
2	standard error	STDERR_FILENO	<i>stderr</i>

- ▶ 4 key system calls for performing file IO
 - ▶ `fd = open(pathname, flags, mode)`
 - ▶ `numread = read(fd, buffer, count)`
 - ▶ `numwritten = write(fd, buffer, count)`
 - ▶ `status = close(fd)`

File I/O – basic method- file copy example



File I/O – open

- ▶ File access mode flags
 - ▶ Read only, Write only and Read-Write
 - ▶ Can retrieve by `fcntl F_GETFL`
- ▶ File Creation flags
- ▶ Open file status flags
 - ▶ Can retrieve and modified by `fcntl F_GETFL`
- ▶ Return file descriptor

```
int open(const char *pathname, int flags,  
... /* mode_t mode */);
```

Flag	Purpose	SUS?
O_RDONLY	Open for reading only	v3
O_WRONLY	Open for writing only	v3
O_RDWR	Open for reading and writing	v3
O_CLOEXEC	Set the close-on-exec flag (since Linux 2.6.23)	v4
O_CREAT	Create file if it doesn't already exist	v3
O_DIRECT	File I/O bypasses buffer cache	
O_DIRECTORY	Fail if <i>pathname</i> is not a directory	v4
O_EXCL	With O_CREAT: create file exclusively	v3
O_LARGEFILE	Used on 32-bit systems to open large files	
O_NOATIME	Don't update file last access time on <i>read()</i> (since Linux 2.6.8)	
O_NOCTTY	Don't let <i>pathname</i> become the controlling terminal	v3
O_NOFOLLOW	Don't dereference symbolic links	v4
O_TRUNC	Truncate existing file to zero length	v3
O_APPEND	Writes are always appended to end of file	v3
O_ASYNC	Generate a signal when I/O is possible	
O_DSYNC	Provide synchronized I/O data integrity (since Linux 2.6.33)	v3
O_NONBLOCK	Open in nonblocking mode	v3
O_SYNC	Make file writes synchronous	v3

File I/O – open (2)

- ▶ Important flags
 - ▶ O_APPEND
 - ▶ Write always append end of file
 - ▶ O_ASYNC
 - ▶ Signal-driven IO, but in Linux this flag no effect, in order to enable this feature, use fcntl flags
 - ▶ O_CLOEXEC
 - ▶ Close the file when calling exec family function
 - ▶ O_CREATE
 - ▶ If file is not exist, file is created
 - ▶ O_DIRECT
 - ▶ Allow File IO to bypass the buffer cache
- ▶ O_EXCL
 - ▶ Conjunction with O_CREATE to indicate that if the file exist, it should not be opened, open fail, errno set to EEXIST
- ▶ O_NONBLOCK
 - ▶ Open file for non-blocking
- ▶ O_SYNC
 - ▶ Open file for synchronous I/O
- ▶ O_TRUNC
 - ▶ If file exist, the content of file is discard and length is truncated to zero

File I/O – open (3)

- ▶ Error from open – strerror(errno)
 - ▶ EACCES
 - ▶ File permission don't allow the calling process to open the file in mode specified by flags
 - ▶ Directory permission don't allow file created or writing
 - ▶ EISDIR
 - ▶ File in pathname is directory and open file for writing
 - ▶ EMFILE
 - ▶ The process resource limit on the number of open file descriptor has been reached
- ▶ ENFILE
 - ▶ System resource limit on the number of open file descriptor has been reached
- ▶ ENOENT
 - ▶ File is not **exist** and **O_CREATE** is not **specified**
- ▶ EROFS
 - ▶ File is read-only but open is for writing
- ▶ ETXTBSY
 - ▶ File is an executable file and it is currently executing. It is not allow to modify a executable file that is running as a process

File I/O - read

► Arguments

- Count specifies the maximum number of bytes to read,
- Buffer is where data placed, the buffer must be at least count bytes long

► Return value

- Return number of bytes actually read
 - Number of byte read is less than count because some reason
 - Interrupt by signal return with errno EINTR
 - New line
- 0 with end of file (EOF), nothing to read
- -1 with error – use strerror(ret) to know what error

```
ssize_t read(int fd, void *buffer, size_t count);
```

```
ssize_t ret;

while (len != 0 && (ret = read (fd, buf, len)) != 0) {
    if (ret == -1) {
        if (errno == EINTR)
            continue;
        perror ("read");
        break;
    }

    len -= ret;
    buf += ret;
}
```

File I/O – read (2)

- ▶ Blocking mode
 - ▶ Read default in blocking mode except regular file
 - ▶ Append zero to end of buffer explicitly
 - ▶ Data can be text, binary interger, C structures in binary form → read don't know what type data comming
 - ▶ Buffer length must be at least one greater than largest string we expect
- ▶ Non blocking mode
 - ▶ Read return immediately with errno set to EAGAIN

```
char buffer[MAX_READ + 1];  
ssize_t numRead;
```

```
numRead = read(STDIN_FILENO, buffer, MAX_READ);  
if (numRead == -1)  
    errExit("read");
```

```
buffer[numRead] = '\0';  
printf("The input data was: %s\n", buffer);
```

File I/O -write

▶ Argument

- ▶ Count is number of bytes write to file descriptor
- ▶ Buffer store the content

▶ Return value

- ▶ Return number of bytes that actually write, this value can be less than count
- ▶ For regular file, a success return from write doesn't guarantee that data has been transfer to disk due to kernel perform buffering of disk I/O to reduce disk activity

```
ssize_t write(int fd, void *buffer, size_t  
count)
```

File I/O - close


- ▶ Free file resource for reuse
- ▶ If process terminate, all descriptor close automatically

```
int close (int fd);
```

File I/O – change file offset - lseek

- ▶ File as sequence of bytes, kernel record file offset
 - ▶ File offset is set to point to the start of file when the file opened
 - ▶ Automatically adjusted by call read or write
 - ▶ Lseek adjust the file offset of open file
- ▶ Argument whence
 - ▶ SEEK_CUR
 - ▶ Offset bytes from the beginning of the file
 - ▶ SEEK_END
 - ▶ Offset bytes from the current file offset
 - ▶ SEEK_SET
 - ▶ The file offset is set to size of the file plus offset → next byte after last byte of file
 - ▶ File hole

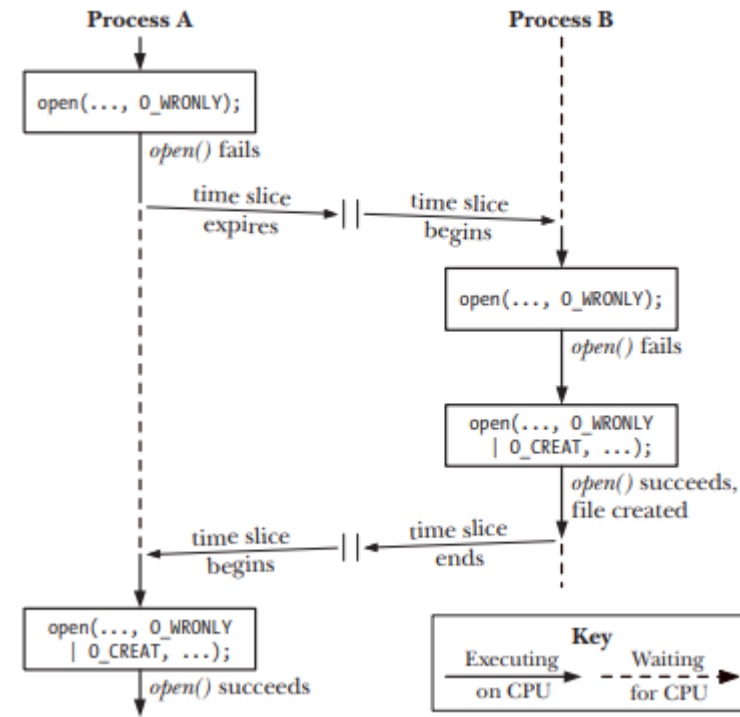
```
off_t lseek (int fd, off_t offset, int whence)
```



```
lseek(fd, 0, SEEK_SET);      /* Start of file */
lseek(fd, 0, SEEK_END);     /* Next byte after the end of the file */
lseek(fd, -1, SEEK_END);    /* Last byte of file */
lseek(fd, -10, SEEK_CUR);   /* Ten bytes prior to current location */
lseek(fd, 10000, SEEK_END); /* 10001 bytes past last byte of file */
```

File I/O – Atomic and Race Conditions

- ▶ All system call are executed atomically
 - ▶ Kernel guarantees that all of the steps in a system call are completed as a single operation without interrupted from other process or threads
 - ▶ Atomic is essential to the successful completion of some operations
 - ▶ Atomic allow avoid race conditions
 - ▶ A situation where the result produced by two processes operate on one shared resources in unexpected way on the relative order in which process gain CPU



File I/O – Appending data to a file

- ▶ Multi-process appending data to the same file to write file
 - ▶ One process seek to end of file
 - ▶ Kernel stop process and allow other running
 - ▶ Other process seek to end of file
 - ▶ Other process write to end of file
 - ▶ Kernel stop other process and allow first process running
 - ▶ First process write to end of file but it replace content of other process
- ▶ Avoid with O_APPEND

```
if (lseek(fd, 0, SEEK_END) == -1)
    errExit("lseek");
if (write(fd, buf, len) != len)
    fatal("Partial/failed write");
```

File I/O –File control operation

- ▶ Control everything relate to file
 - ▶ Duplicate file descriptor
 - ▶ File descriptor flags
 - ▶ File status flags
 - ▶ File locking
 - ▶ IO availability signal
 - ▶ File change notification

```
int fcntl(int fd, int cmd, ...);
```


FILE I/O – File status flags

► fcntl

- Retrieve or modify the access mode and open file status flags
- F_GETFL command
- F_SETFL command
 - Modify O_APPEND, O_NONBLOCK, O_NOATIME, O_ASYNC, O_DIRECT

```
int flags, accessMode;

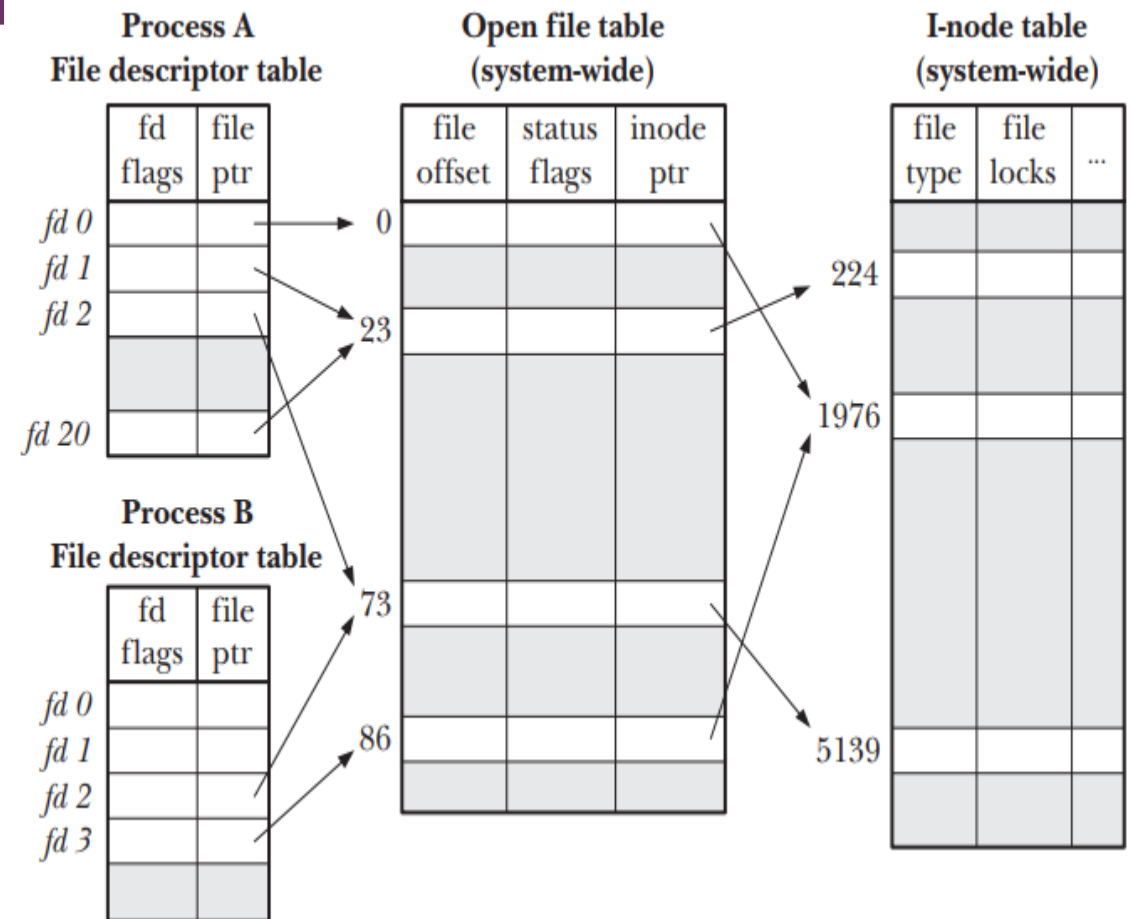
flags = fcntl(fd, F_GETFL);
if (flags == -1)
    errExit("fcntl");

if (flags & O_SYNC)
    printf("writes are synchronized\n");
accessMode = flags & O_ACCMODE;
if (accessMode == O_WRONLY || accessMode == O_RDWR)
    printf("file is writable\n");

flags = fcntl(fd, F_GETFL);
if (flags == -1)
    errExit("fcntl");
flags |= O_APPEND;
if (fcntl(fd, F_SETFL, flags) == -1)
    errExit("fcntl");
```

File I/O – Relationship between File Descriptor and Open File

- ▶ Three data structures that maintain by kernel
 - ▶ Per-process file descriptor
 - ▶ set of flags controlling the operation of the file descriptor (close on exec)
 - ▶ The system wide table of open file descriptor
 - ▶ the current file offset
 - ▶ status flags specified when opening the file
 - ▶ O_APPEND, O_NONBLOCK, O_ASYNC
 - ▶ the file access mode(readonly, writeonly or read write)
 - ▶ settings relating to signal-driven I/O
 - ▶ a reference to the i-node object
 - ▶ File system inode table
 - ▶ file type (e.g., regular file, socket, or FIFO) and permissions
 - ▶ a pointer to a list of locks held on this file
 - ▶ various properties of the file, including its size and timestamps



File I/O – Duplicate file descriptor

- ▶ I/O redirection
 - ▶ Standard error redirected to the same place with standard output
 - ▶ Sent both standard output and standard error to results.log
- ▶ dup
 - ▶ Take an open file descriptor and return a new descriptor that refer to the same open file descriptor
- ▶ dup2
 - ▶ Make a duplicate of file descriptor old to new
 - ▶ If newfd is opened, it closed

```
$ ./myscript > results.log 2>&1
```

```
int dup(int oldfd);
```

```
close(2);           /* Frees file descriptor 2 */  
newfd = dup(1);      /* Should reuse file descriptor 2 */
```

```
int dup2(int oldfd, int newfd);
```

File I/O at specified Offset

- ▶ Pread and pwrite
 - ▶ Operate like read and write
 - ▶ The file IO is performed at location specified by offset than current file offset
 - ▶ The file offset left unchanged
- ▶ Call pread equivalent call lseek + read but atomically
 - ▶ These function is useful in multithread applications
 - ▶ All open file descriptor shared by all threads
 - ▶ File offset of each file is global to all threads
 - ▶ Avoid race condition when read and write simultaneously

```
ssize_t pread(int fd, void *buf, size_t count, off_t offset);
```

Returns number of bytes read, 0 on EOF, or -1 on error

```
ssize_t pwrite(int fd, const void *buf, size_t count, off_t offset);
```

Returns number of bytes written, or -1 on error

```
off_t orig;
```

```
orig = lseek(fd, 0, SEEK_CUR);    /* Save current offset */  
lseek(fd, offset, SEEK_SET);  
s = read(fd, buf, len);  
lseek(fd, orig, SEEK_SET);       /* Restore original file offset */
```

File I/O – scatter/gather IO

► Scatter/Gather IO

- Transfer multiple buffer of data in single system call
- Set of buffers to be transferred is defined by the array `iov`

► Scatter input

- `readv` perform scatter input
- read a contiguous sequence of bytes from the file descriptor to buffer's `iovec`
- `readv` is atomically → when reading from a file, the range of byte in `iov`'s buffer is continuous even if another process sharing the same file offset attempts to manipulate the offset at the same time
- `Readv` return number of byte read or 0 if end-of-file was encountered
- If insufficient data, some of last is partially filled

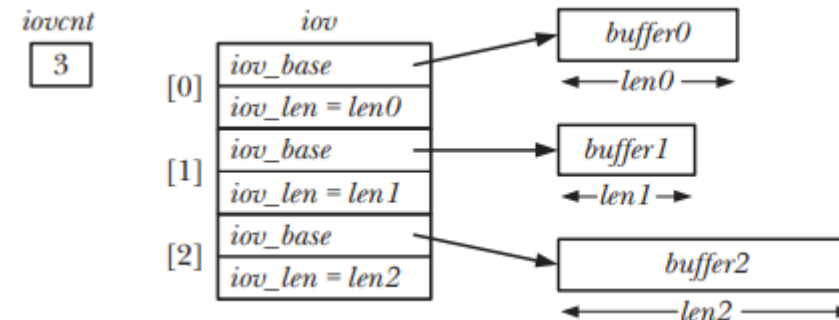
```
ssize_t readv(int fd, const struct iovec *iov, int iovcnt);
```

Returns number of bytes read, 0 on EOF, or -1 on error

```
ssize_t writev(int fd, const struct iovec *iov, int iovcnt);
```

Returns number of bytes written, or -1 on error

```
struct iovec {  
    void *iov_base;    /* Start address of buffer */  
    size_t iov_len;    /* Number of bytes to transfer to/from buffer */  
};
```



File I/O – scatter/gather IO (2)

- ▶ Gather output
 - ▶ Writev system call perform gather output
 - ▶ Write as sequence of continuous bytes
 - ▶ Buffer gather in array order
 - ▶ Writev is atomically → all requested data is written continuously to the file
 - ▶ Like write, writev can return partially write. Let check the result to indicate
 - ▶ Readv and writev is convenience and speed
 - ▶ Writev can implement as copy of user buffers and call write
 - ▶ Call write multiple time but two way inconvenience and slow

- ▶ Performing scatter-gather I/O at a specified offset

```
ssize_t preadv(int fd, const struct iovec *iov, int iovcnt, off_t offset);
```

Returns number of bytes read, 0 on EOF, or -1 on error

```
ssize_t pwritev(int fd, const struct iovec *iov, int iovcnt, off_t offset);
```

Returns number of bytes written, or -1 on error

FILE I/O – truncating the file

- ▶ Resize the file length
 - ▶ If file length is longer than length, excess data is lost
 - ▶ If file length is shorter than length, file will padding with sequence of null byte
 - ▶ The command use to discard the content of file

```
int truncate(const char *pathname, off_t length);  
int ftruncate(int fd, off_t length);
```

File I/O- creating temporary file

- ▶ Create temporary file when program is running
 - ▶ Create a file in /tmp directory and return file descriptor
 - ▶ Template argument takes the form of a pathname in which the last 6 characters must be XXXXXX
 - ▶ 6 XXXXXX will be replaced with a string that makes the filename unique
 - ▶ Template is modified → it must be specified as a character array rather than string constant

```
int mkstemp(char *template);
```


Kernel Buffering of File I/O

- ▶ Working with disk
 - ▶ Read/Write system call don't directly initiate disk access
 - ▶ They copy data between a user-space buffer and a buffer in kernel buffer cache
 - ▶ Write(fd, "abc", 3)
 - ▶ Write return immediately
 - ▶ Some later point, kernel writes its buffer to disk
 - ▶ System call is not synchronized with disk
 - ▶ If another process attempts to read these bytes of file, the kernel supplies data from buffer cache
- ▶ Read(fd, buf, 3)
 - ▶ Kernel read data from the buffer until is exhausted , kernel reads next segment of the file into the buffer cache
 - ▶ This design is allow read and write to be fast, process don't need to wait on slow disk, and reduce the number of disk transfer that kernel perform
- ▶ Two way to control kernel buffer of File I/O
 - ▶ Use system call

```
int fsync(int fd);  
void sync(void);
```
 - ▶ Use O_SYNC when open file

stdio library – standard IO

- ▶ File streams
- ▶ Standard IO

File descriptor	Purpose	POSIX name	<i>stdio</i> stream
0	standard input	STDIN_FILENO	<i>stdin</i>
1	standard output	STDOUT_FILENO	<i>stdout</i>
2	standard error	STDERR_FILENO	<i>stderr</i>

- ▶ Relationship between stdio and File IO

stdio library – API – open the file

*FILE * **fopen** (const char *filename, const char *opentype)*

<i>type</i>	Description	open(2) Flags
r or rb w or wb a or ab	open for reading truncate to 0 length or create for writing append; open for writing at end of file, or create for writing	O_RDONLY O_WRONLY O_CREAT O_TRUNC O_WRONLY O_CREAT O_APPEND
r+ or r+b or rb+ w+ or w+b or wb+	open for reading and writing truncate to 0 length or create for reading and writing	O_RDWR O_RDWR O_CREAT O_TRUNC
a+ or a+b or ab+	open or create for reading and writing at end of file	O_RDWR O_CREAT O_APPEND

Restriction	r	w	a	r+	w+	a+
file must already exist	•			•		
previous contents of file discarded		•			•	
stream can be read	•			•	•	•
stream can be written		•	•	•	•	•
stream can be written only at end			•			•

stdio library – API – close the file

```
int fclose (FILE *stream)
```

stdio library – API – read/write the file

fread (*void *data, size_t size, size_t count, FILE *stream*)

size_t **fwrite** (*const void *data, size_t size,
size_t count, FILE *stream*)

stdio library – API – file positioning

*long int **ftell** (FILE *stream)*

*int **fseek** (FILE *stream, long int offset,
int whence)*

Buffering in the stdio library

- ▶ Buffering of data into large block
 - ▶ Reduce system call and overhead
 - ▶ Increase performance when operate on disk
- ▶ Setting the buffering mode of stdio stream
 - ▶ If buf is NULL, library automatically allocates a buffer for use with stream at least BUFSIZ 512 byte
 - ▶ Mode
 - ▶ `_IONBF` Don't buffer IO, call system call read/write immediately, buf and size is ignored
 - ▶ `_IOLBF` Line buffer, data is buffer until a newline character is output, for input, data is read a line at a time
 - ▶ `_IOFBF` Fully buffer I/O is default mode
- ▶ Flushing a stdio buffer

```
int setvbuf(FILE *stream, char *buf, int mode, size_t size);
```

```
int fflush(FILE *stream);
```


stdio library – API -buffering

- ▶ Performance reason
- ▶ Set buffer for buffering
 - ▶ Atleast BUFSIZ – 512 bytes
 - ▶ Set NULL to disable buffering

```
void setbuf(FILE *restrict fp, char *restrict buf );
```

Review

- ▶ strace
- ▶ fallocation -l 100M /tmp/bigfile
- ▶ time