



Tin học Nhà Uyên

PROCESS IN LINUX

TIN HOC NHA UYEN

CONTENT

- PROCESS ID
- PROCESS MEMORY LAYOUT
- VIRTUAL MEMORY MENAGEMENT
- STACK AND STACK FRAME
- COMMAND LINE ARGUMENT
- ENVIRONMENT VARIABLES
- PROCESS CREATION
- PROCESS TERMINATION
- MONITORING CHILD PROCESS
- PROGRAM EXECUTION

PROCESS ID AND PARENT PROCESS ID

- Process ID
 - Uniquely identifies the process on the system
 - Used by varieties of system call like kill
 - Limited to 32767
- Parent Process ID
 - Process create child process
 - pstree

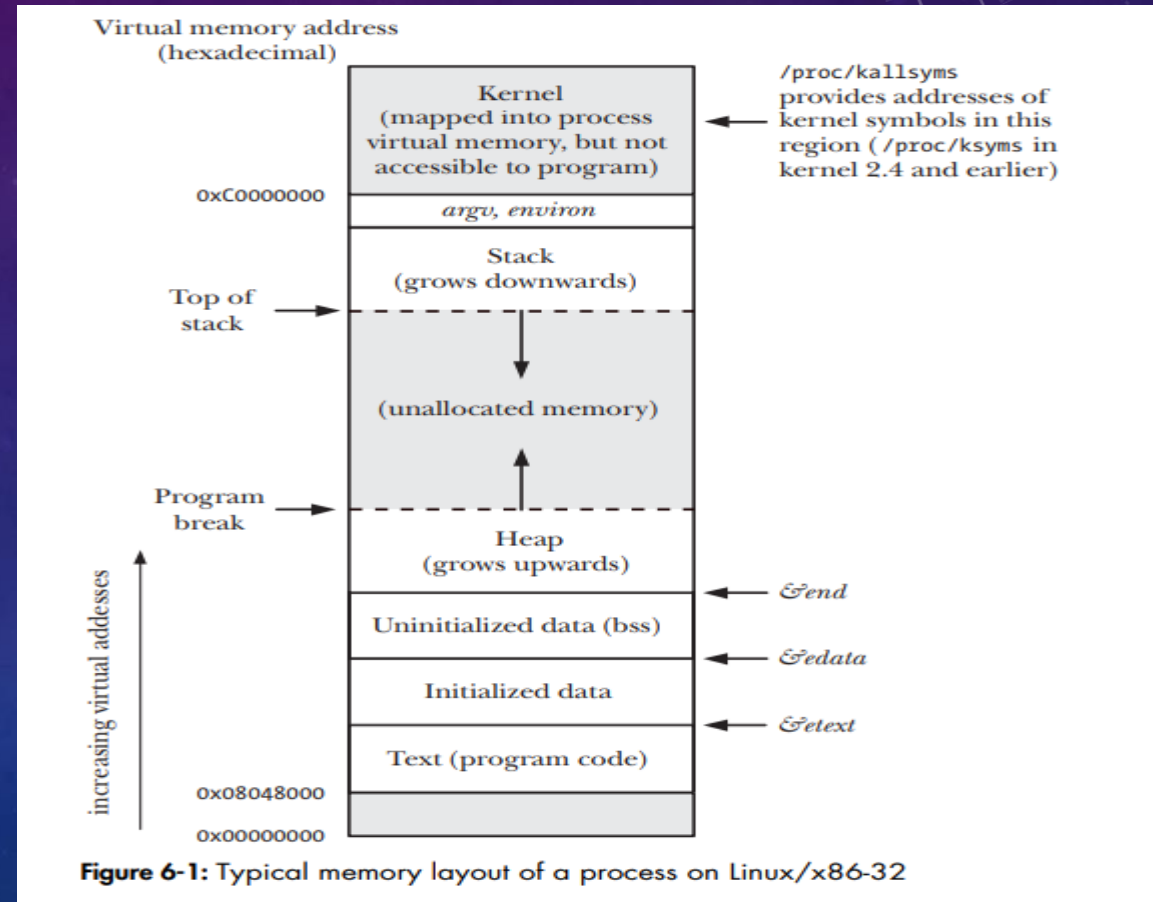
```
pid_t getpid(void);
```

```
pid_t getppid(void);
```

```
pstree command
```

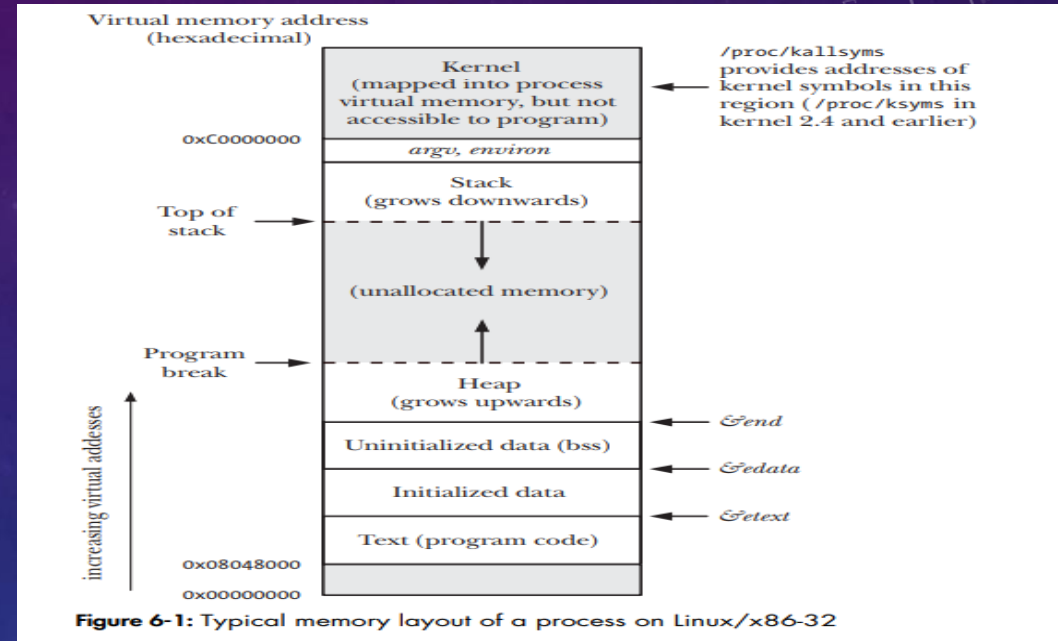
PROCESS MEMORY LAYOUT

- Text segment
 - Machine-language instructions of program run by the process
 - Read-only area that process don't accidentally modify it own intructions via bad pointer access
 - When multiple process run the same program, text segment shareable so that a single copy of program code can be mapped into virtual address of all of processes
- Initialized data segment
 - Global and static variables that explicitly initialized
- Uninitialized data segment
 - Global and static variables that explicitly initialized



PROCESS MEMORY LAYOUT

- Stack
 - Dynamically growing and shrinking segment containing stack frame
 - One stack frame is allocated for each currently called function
 - One frame store local variables (automatic variables), argument and return value
- Heap
 - Dynamically allocation by run time
 - Size command



```
thientran@ubuntu:~/Desktop/training/process$ size bin_stack
text    data    bss     dec     hex filename
1198    292 1089568 1091058 10a5f2 bin_stack
```

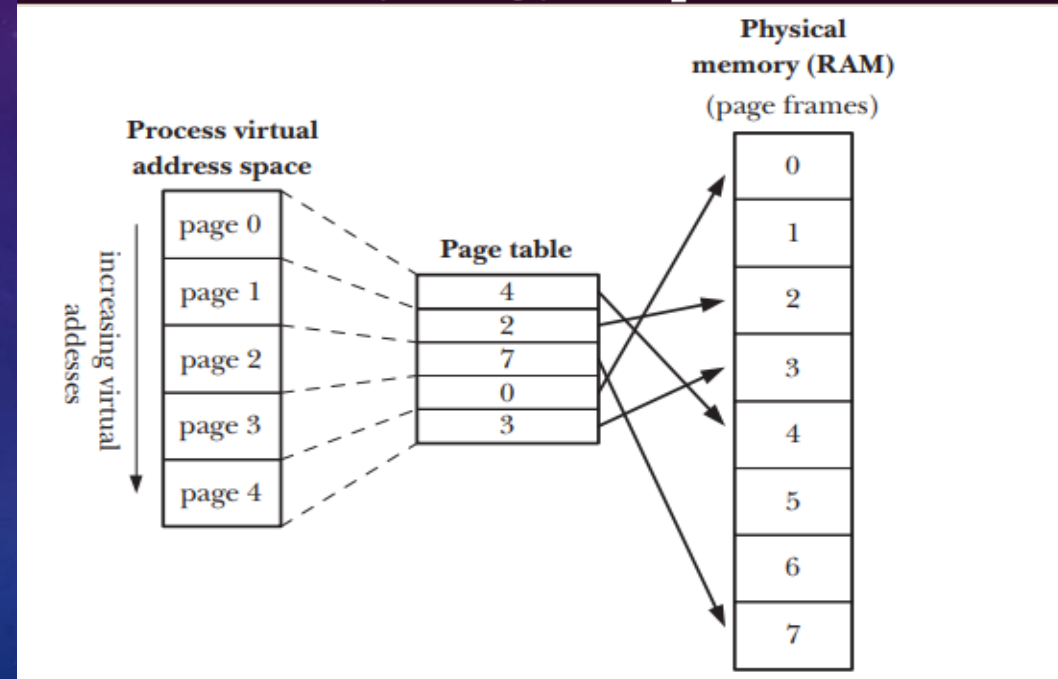
VIRTUAL MEMORY MANAGEMENT

- To make efficient use of both the CPU and RAM by exploiting a typical property of most programs : Locality of reference
- Two kinds of locality
 - Spatial locality
 - The tendency of program to reference memory addresses that recently access
 - Temporal locality
 - The tendency of program to reference memory addresses that access recent past (due to loop)

VIRTUAL MEMORY MANAGEMENT

- Virtual memory split memory used by each program into small fixed size unit called pages
- Only some of the pages of a program need to be resident in physical memory page frames, these pages are called by resident-set
- A reserved-area of disk space used to supplement the computer RAM is called swap-area
- When a process references a page that is not currently resident in physical memory, a page fault occurs, kernel suspends execution of process while the page is loaded from disk into memory

```
thientran@ubuntu:~/Desktop/training/process$ getconf PAGE_SIZE
4096
thientran@ubuntu:~/Desktop/training/process$ getconf PAGE_SIZE
4096
thientran@ubuntu:~/Desktop/training/process$
```



VIRTUAL MEMORY MANAGEMENT

- Page table
 - Describe location of each page in process virtual address space
 - Each entry in page table indicate it is currently reside on RAM or Disk
 - Not all address ranges in process virtual address virtual space require page-table entries, it is not necessary to maintain corresponding page-table entries
 - If process access in a address that is not corresponding page table entry. It receive a SIGSEGV

```
1 #include <stdio.h>
2 #include <stdlib.h>
3
4 int main(int argc, char *argv[])
5 {
6     char buf[5];
7
8     buf[0] = 0x45;
9     buf[2] = 0x45;
10    buf[5000] = 0x49;
11
12    printf("%x ", buf[5000]);
13    exit(EXIT_SUCCESS);
14 }
```


VIRTUAL MEMORY MANAGEMENT

- Process range of valid virtual addresses
 - Stack grows downward beyond limits
 - Memory allocated or deallocated on the heap
 - malloc or free
 - Shared memory attach or detached
 - shmat or shmdt
 - Memory mapping mapped or unmapped
 - mmap or munmap

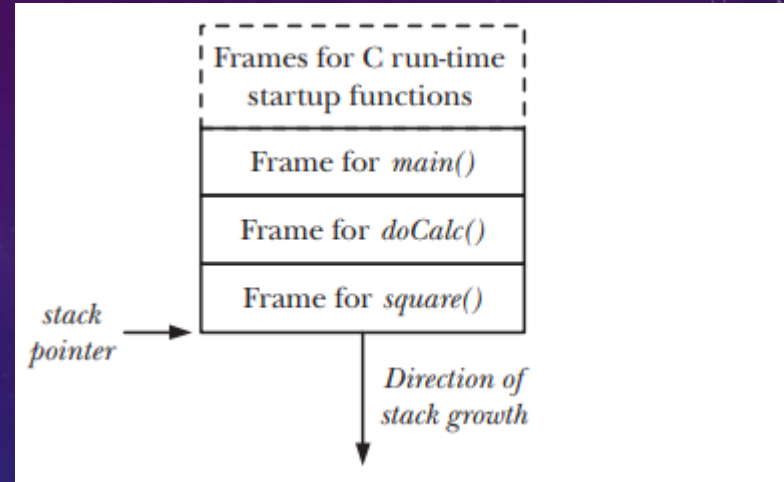
VIRTUAL MEMORY MANAGEMENT

- Advantages

- Process isolated from other
 - One process cannot read or write memory of other process or kernel
 - Accomplished by having the page-table entries for each process point to distinct sets of physical pages in RAM
- When appropriate, two or more process can share memory
 - Multi-process execute the same program can share single copy of program code as multiple process run the same binary or load the same share library
 - IPC to communication between process
 - Accomplished by having page-tables entries in difference process refer to same page of RAM
- Easy to implement memory protection
 - by allow some processes can readable, writeable or executable on some page-table entries
- Programmer don't need to concern about physical layout of program
- All part of program reside in memory then program run faster and memory capacity exceed physical memory

STACK AND STACK FRAME

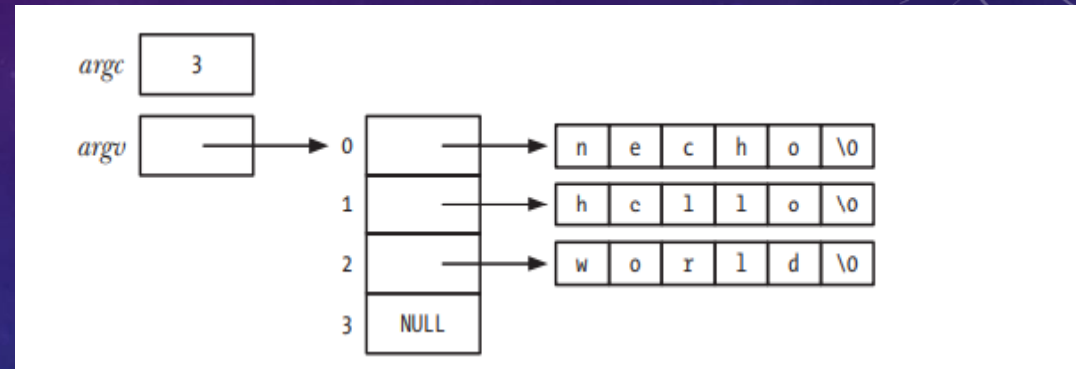
- Stack pointer
 - Special purpose register track the current top of stack
 - Each time a function is called , an additional frame is allocated on stack, this frame is remove when function returns
- User stack
 - Include function argument and local variables
 - Automatic variable vs static variable
 - Call linkage information
 - Each function use certain CPU register (PC, EAX), each a function call another , a copy of these registers saved in the called function stack frame so that function return
 - A function call itself recursively, it has multiple stack frame



```
(gdb) bt
#0  0x00110422 in __kernel_vsyscall ()
#1  0x00894651 in raise () from /lib/tls/i686/cmov/libc.so.6
#2  0x00897a82 in abort () from /lib/tls/i686/cmov/libc.so.6
#3  0x08048469 in square (x=9973) at stack.c:11
#4  0x0804847a in docalc (val=9973) at stack.c:22
#5  0x080484dc in main (argc=1, argv=0xbfae0084) at stack.c:44
(gdb) █
```

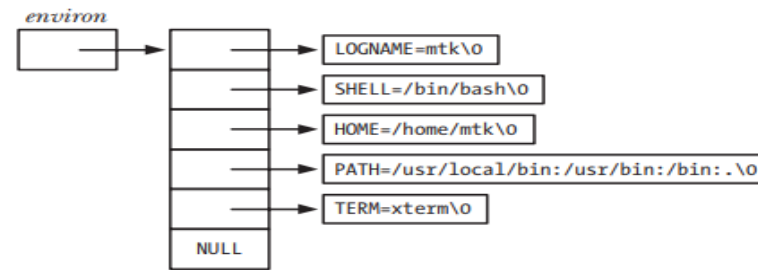
COMMAND LINE ARGUMENT

- `main(int argc, char *argv[])`
- `argc`
 - How many command line argument there are
- `argv`
 - `argv[0]` is name of program
 - convert argument to number by `strtol` or `atoi`
 - use `getopt` library
- See command line argument of a process
 - `/proc/PID/cmdline`



ENVIRONMENT VARIABLES

- Environment list
 - Array of strings associated for process
 - When a process created, it inherit a copy of parent environment
 - A common use of environment variables is in the shell. By placing values in shell environment, the shell ensure it will pass to process that execute in the shell
- export command
 - export
 - export PATH=\$PATH:/path_to_dir
 - bashrc
- PATH and LD_LIBRARY_PATH environment
- environ global variables
- APIs



```
#include <stdio.h>
#include <stdlib.h>

extern char **environ;
int main(int argc, char *argv[])
{
    char **ep;
    for (ep = environ; *ep != NULL; ep++)
        puts(*ep);
    exit(EXIT_SUCCESS);
}
```

```
char *getenv(const char *name);
```

```
int setenv(const char *name, const char *value, int overwrite);
```

```
int unsetenv(const char *name);
```

```
int clearenv(void)
```

PROCESS CREATION- OVERVIEW OF BASIC SYSTEM CALL

- `fork()`
 - Allow parent process create child process
 - Child obtains copies of the parent's stack, data, heap and text segments
- `exit(status)`
 - Terminate a process, making all resources like memory, open file descriptor used by process now available for kernel reallocation
 - status determines the termination status for process
- `wait(status)`
 - If process is not terminated , wait will suspends execution of the process until one of children has terminated
 - If process terminated, wait will return immediately with child termination status
- `execve(pathname, argv, envp)`
 - Loads program into process memory with argument argv and environment list envp
 - The existing program text is discarded, the stack, data, heap segments are freshly created

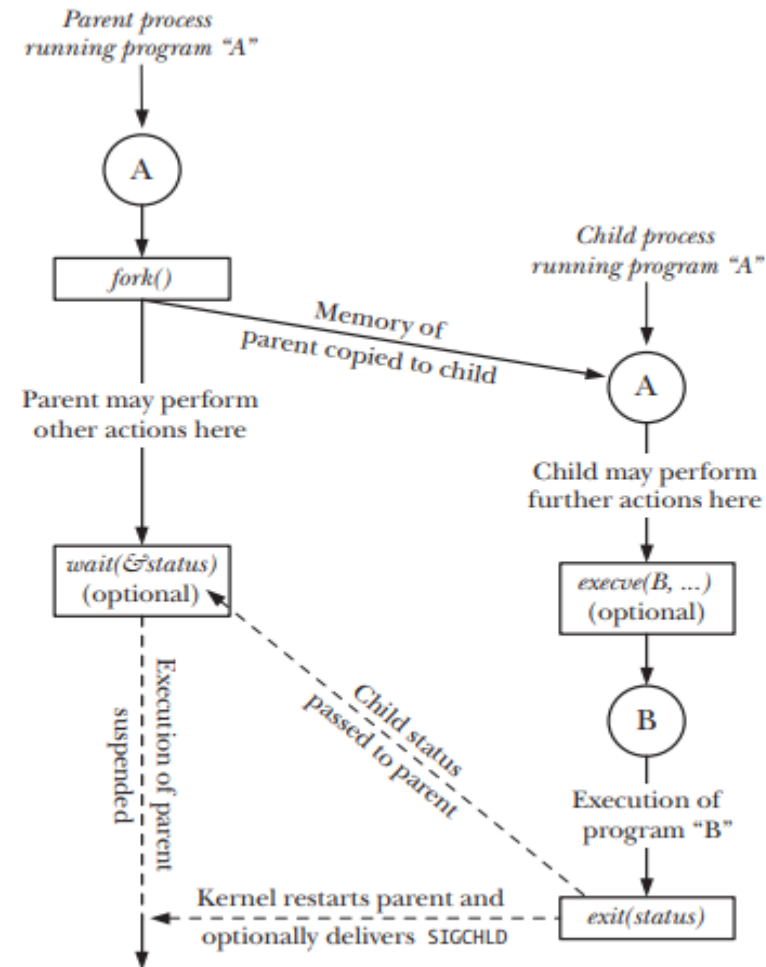


Figure 24-1: Overview of the use of `fork()`, `exit()`, `wait()`, and `execve()`

CREATING NEW PROCESS WITH FORK

- Creating new process is useful way of dividing tasks as network server receive request from multi-clients
- fork creates a new process, the child, that is an almost exact duplicate of calling process, the parent. When fork completed, two process exist and execution of two process is continuing
- Two processes are executing the same program text but they have separate copies of stack, data and heap segment
- The child's stack data and heap segment are initially exact duplicates of corresponding part of parent memory.
- After fork, each process can modify the variables in stack, data, heap segment without affecting the other process
- fork will return child process id to parent process, 0 to child process and -1 if something error

```
pid_t fork(void);
```

In parent: returns process ID of child on success, or -1 on error;
in successfully created child: always returns 0

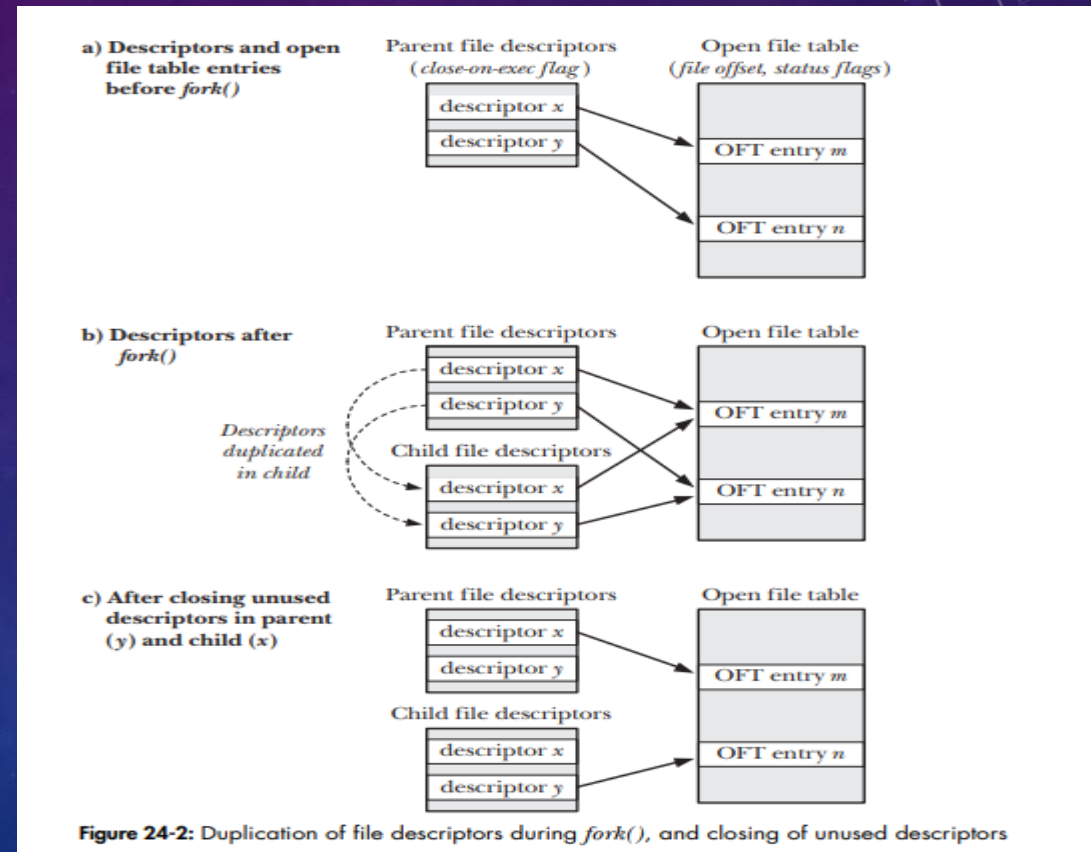
```
pid_t childPid;          /* Used in parent after successful fork()
                           to record PID of child */
switch (childPid = fork()) {
case -1:                  /* fork() failed */
    /* Handle error */

case 0:                   /* Child of successful fork() comes here */
    /* Perform actions specific to child */

default:                  /* Parent comes here after successful fork() */
    /* Perform actions specific to parent */
}
```

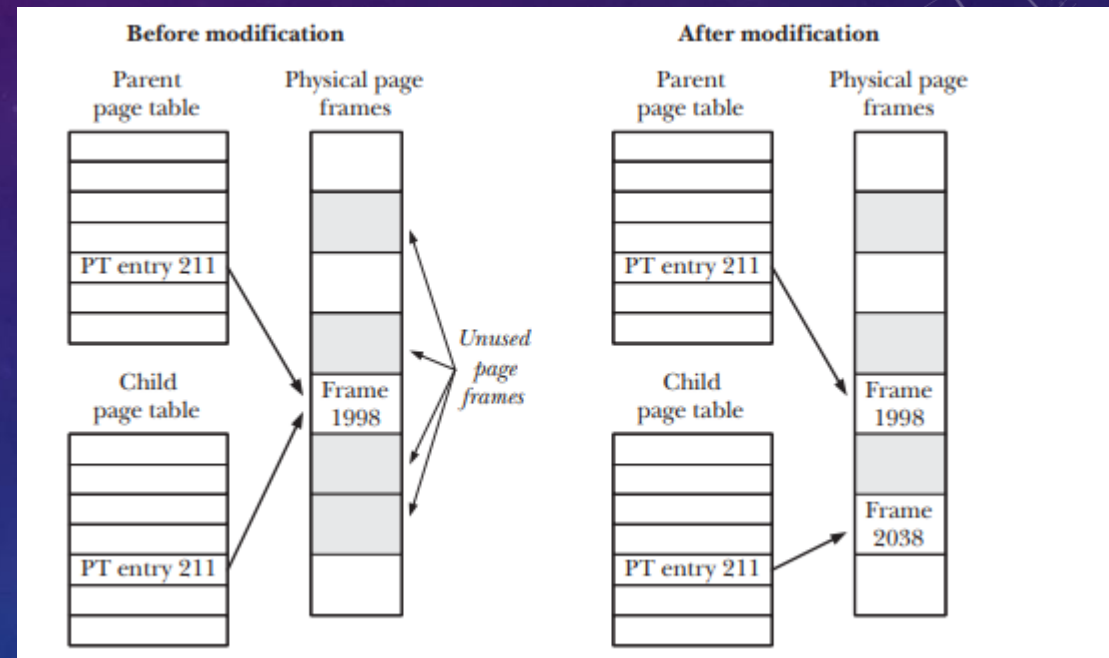

FILE SHARING BETWEEN PARENT AND CHILD

- After fork, the child receives duplicate of all of the parent file descriptor, it mean parent and child refer to the same open file descriptor
- All attributes of open file is shared between the parent and child
- For example, the child modify file offset, this change is visible through the corresponding descriptor in parent
- If the parent and the child both writing the file, sharing offset ensure that the two processes don't override each other's output
- However it is not prevent the output of two processes from being randomly intermingled
- The code should close unused descriptor



MEMORY SEMATIC OF FORK

- The wasteful copy of parent text, data, heap and stack
 - Especially after fork is exec
 - Exec will discard current text, data, heap, stack and replace by its own program
- Kernel use some techniques to avoid such wasteful copying
 - With text segment, this area is read-only, so that process cannot modify its own text segment, then
 - All entries in per-process per-tables in child refer to same virtual memory page frame already used by parent
 - For pages in data, heap and stack, kernel employs copy-on-write technique

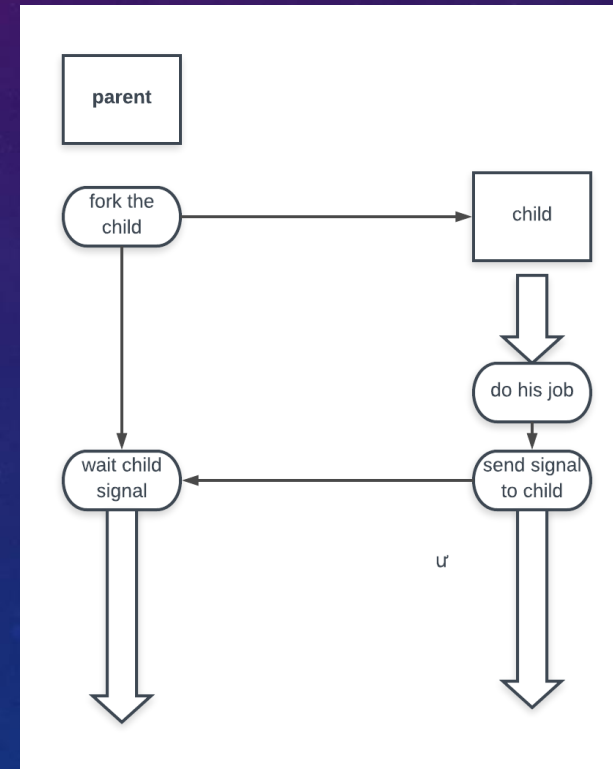


RACE CONDITIONS AFTER FORK

- Parent scheduled first in almost times but not guarantee
- `/proc/sys/kernel/sched_child_runs_first` work?
- Several way to guarantee the order of process of parent and child
 - Semaphore
 - File lock
 - Messaging through IPC
 - Signal

AVOID RACE CONDITION BY SYNCHRONIZING WITH SIGNAL

- Parent wait his child finish the job and process



PROCESS TERMINATION – TERMINATING A PROCESS WITH EXIT AND _EXIT

```
void _exit(int status);
```


- Exit a process
 - Two way for terminating a process
 - Normal way through exit or program finish
 - Abnormal way through receiving a signal
 - Actions are performed by calling exit
 - Exit handler
 - stdio buffer stream flush
 - _exit system call involved
- Process termination detail
 - Open file descriptor, directory streams ... are closed
 - File lock held by process are released
 - Attached share memory are detached
 - Semaphores are closed
 - Message queue are closed
 - Memory lock is removed
 - Memory mapping is unmapped

EXIT HANDLER – ATEXTIT/ON_EXIT

```
int atexit(void (*func)(void));
```

- Cleanup action

- Some libraries or application need to perform cleanup action automatically when process exit
- Linux provide a programmer supplied function that is registered in lifetime of process and automatically called during normal process terminate by call exit
- Exit handler is not called if a program call `_exit()` system call directly or if process terminated abnormal by a signal



func no have argument and return value

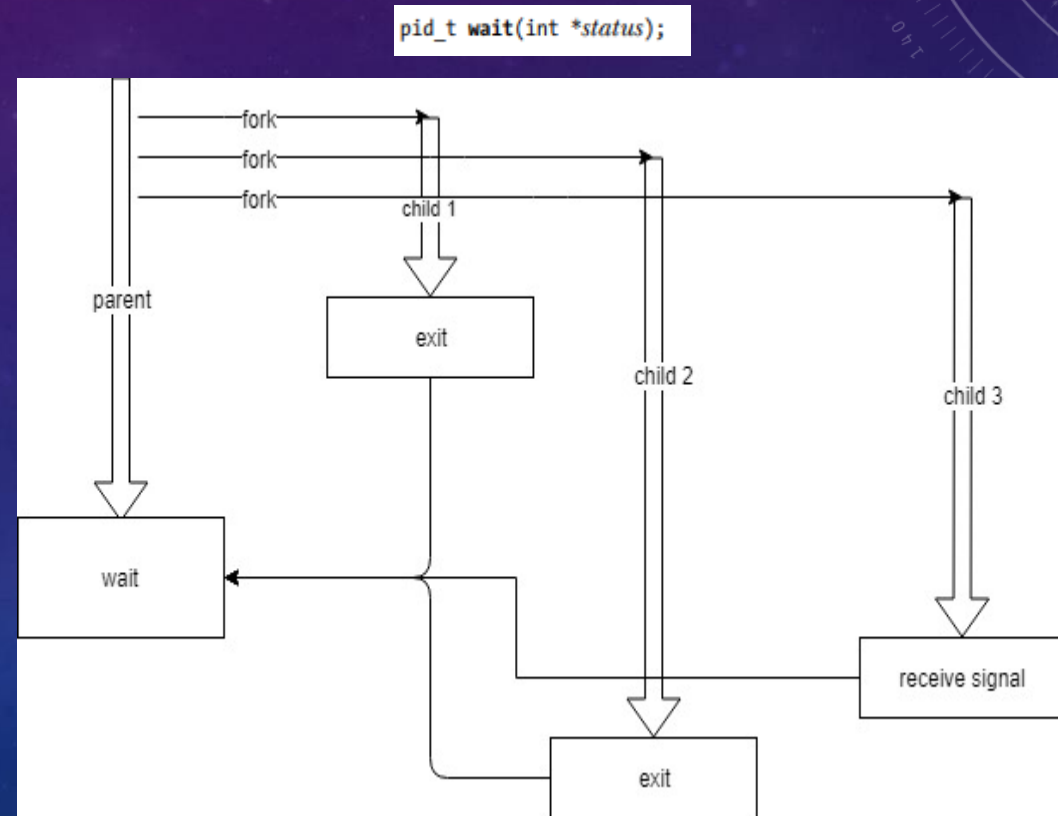
- When a program invoke exit, these func are called in reserve order of registration
- If one of exit handler call `_exit` or raise a signal, process will terminate and ignore the remaining exit handler
- A child process inherit parent exit handler registration after fork
- When a process perform exec , all exit handler registration is remove

INTERACTION BETWEEN FORK, STDIO BUFFER AND EXIT

- Duplicate output when redirect output
 - Flush stdout before fork
 - Call setbuf to disable buffering on stdio stream
 - One process call exit, other call `_exit` system call

MONITORING CHILD PROCESS – WAIT SYSTEM CALL

- Waiting on child process
 - wait system call
 - Wait one of the children of calling process terminate and return termination status to status argument
 - If no child of calling process has yet terminated, the call blocks until one of children terminate
 - If one of children terminate already, wait return immediately
 - If status is not NULL, information about how the child terminated is returned in the status argument
 - The kernel adds the process CPU times and resource usage statistic to running totals for all children
 - Wait return the process id of children process that has terminated



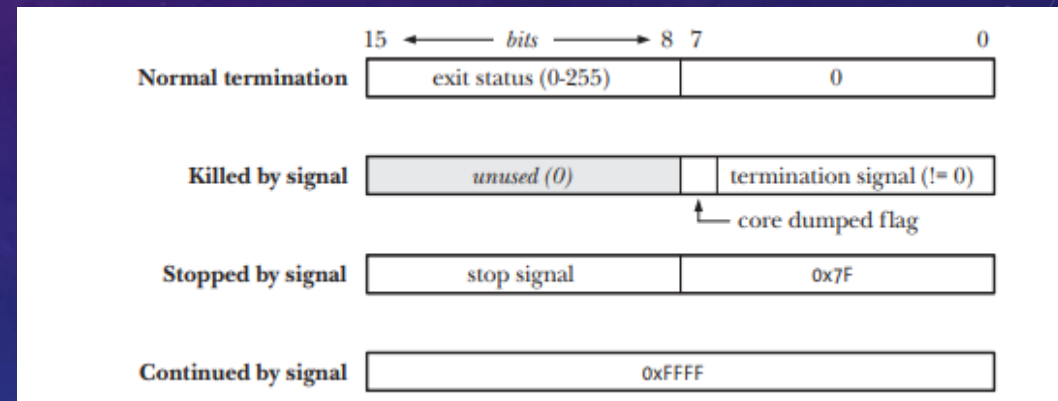
WAITPID SYSTEM CALL

```
pid_t waitpid(pid_t pid, int *status, int options);
```

- Limitations of wait
 - Parent cannot wait a specific children
 - If no child has yet terminate, wait always block, sometime, program need parent do some other work
 - Parent don't know when child receive a SIGSTOP
- pid argument enables the selection of the child to be waited for
 - If $pid > 0$, wait process has process id is pid
 - If $pid = 0$, wait any child
 - Status argument as wait
- Options
 - WUNTRACED return when child receive a SIGSTOP signal
 - WCONTINUED return when a stop process is resume
 - WNOHANG return immediately if children that has process id is pid no have state change (poll)

VALUE OF STATUS

- Status value returned by wait and waitpid allow distinguish following event of the child
 - The child terminating by calling exit
 - The child terminated by the delivery of an unhandled signal
 - The child stop by signal SIGSTOP
 - The child resume by signal SIGCONT
- Macro
 - WIFEXITED(status) – normal exit
 - WIFSIGNALED(status) –unhandled signal
 - WIFSTOPPED(status) – signal stop
 - WIFCONTINUED(status) –signal continue

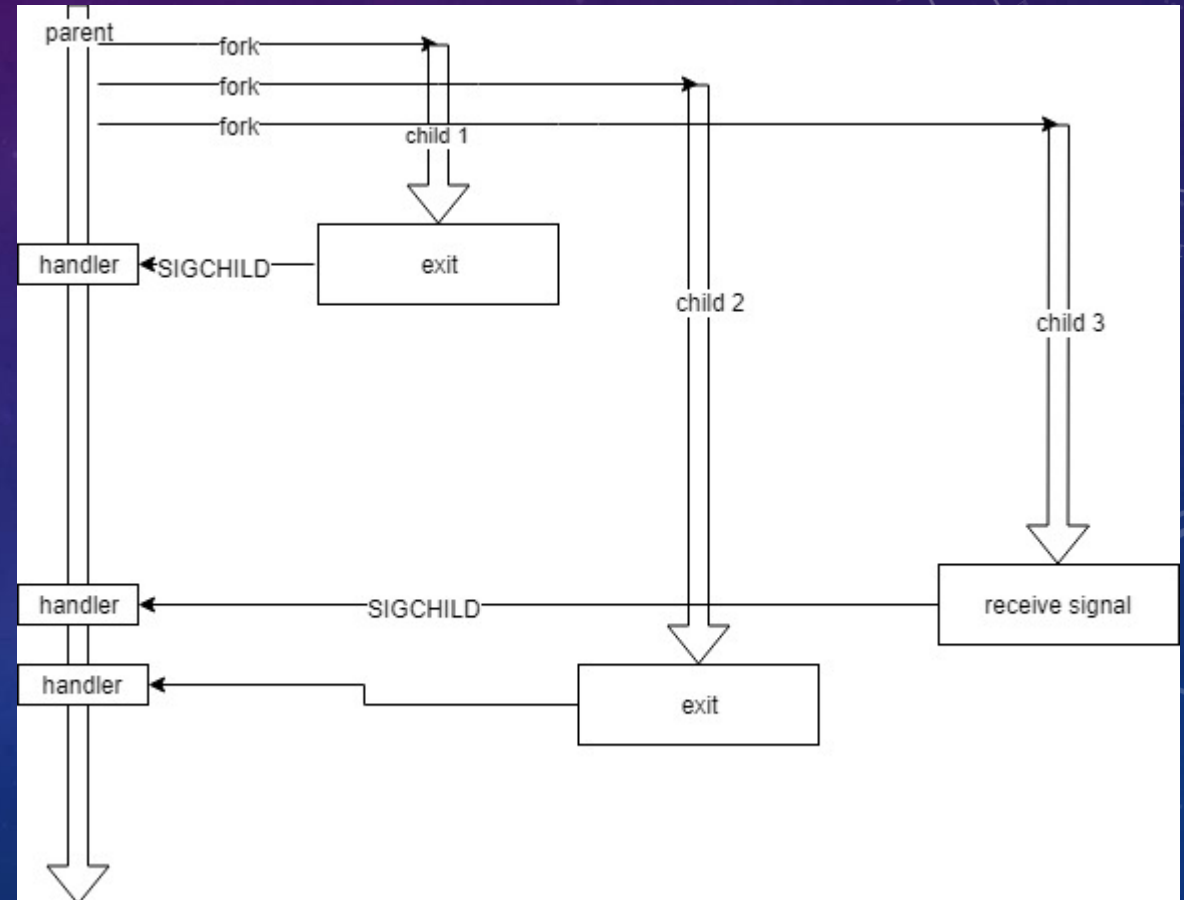


ORPHANS AND ZOMBIES

- Who become an parent of orphaned child?
 - init process –pid 1
 - getppid function
- What happen if a children terminate before its parent call wait
 - Parent be permitted call wait even when children terminated, kernel maintain terminated children at zombie state
 - Almost children resource already release back and reallocate for other process, only info in kernel process table entry that include child pid, termination status and resource usage status
 - When parent call wait, this entry removed from table
 - Otherwise, this entry exist forever until parent process terminated

SIGCHILD SIGNAL

- Child process terminated
 - An asynchronously event that parent is not predict
 - Parent can call wait and waitpid in blocking (without WNOHANG)
 - Parent can call waitpid in nonblocking (with WNOHANG)
 - Both blocking or non-blocking is inconvenient. Then child send SIGCHILD to parent each time it terminate



PROGRAM EXECUTION

- Executing a program with `execve` system call
 - Load a program into a process memory
 - Old program is discarded and process's stack, data, heap, text segment are replaced by new program
 - Most frequent use of `execve` is in the child produced by a fork

```
int execve(const char *pathname, char *const argv[], char *const envp[]);
```

- Pathname : name of binary
- argv : command line path to bin
- envp: environment list

FILE DESCRIPTOR AND EXEC

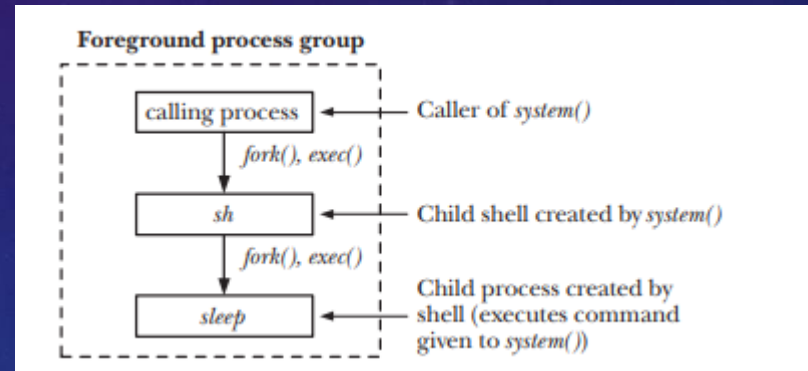
- Open file descriptor through exec
 - These files remain open across the exec and available for new program
 - Shell take advance of this feature to handle I/O redirection for program that it execute.
 - i.e. ls /tmp > dir.txt
 - Fork the child, pass the command
 - Child open dir.txt so that it has file descriptor 1
 - Close stdout, the open file descriptor of dir.txt
 - Dup2 fd of dir.txt and stdout
 - Exec the ls command, ls write to stdout which currently is dir.txt

```
thientran@ubuntu:~/Desktop/training/process$ ls /tmp/ > dir.txt
thientran@ubuntu:~/Desktop/training/process$ cat dir.txt
gedit.thientran.660062292
keyring-sgPMch
orbit-gdm
orbit-thientran
pulse-Bjz5ASS0YxiT
pulse-PKdhtXMmr18n
ssh-XBNQXF1738
virtual-thientran.wxEBI6
VMwareDnD
vmware-root
vmware-root-826451963
vmware-thientran
```

EXECUTING IN A SHELL COMMAND

```
int system(const char *command);
```

- system function
 - system("ls /dev/ | grep rtc") or system("sleep")
 - System function create a child process that involve a shell to execute command



PROCESS FUNCTION REVIEW

- GETPID/GETPPID
- FORK
- EXIT
 - ATEXT/ON_EXIT
- WAIT/WAITPID
- EXECVE
- SYSTEM