



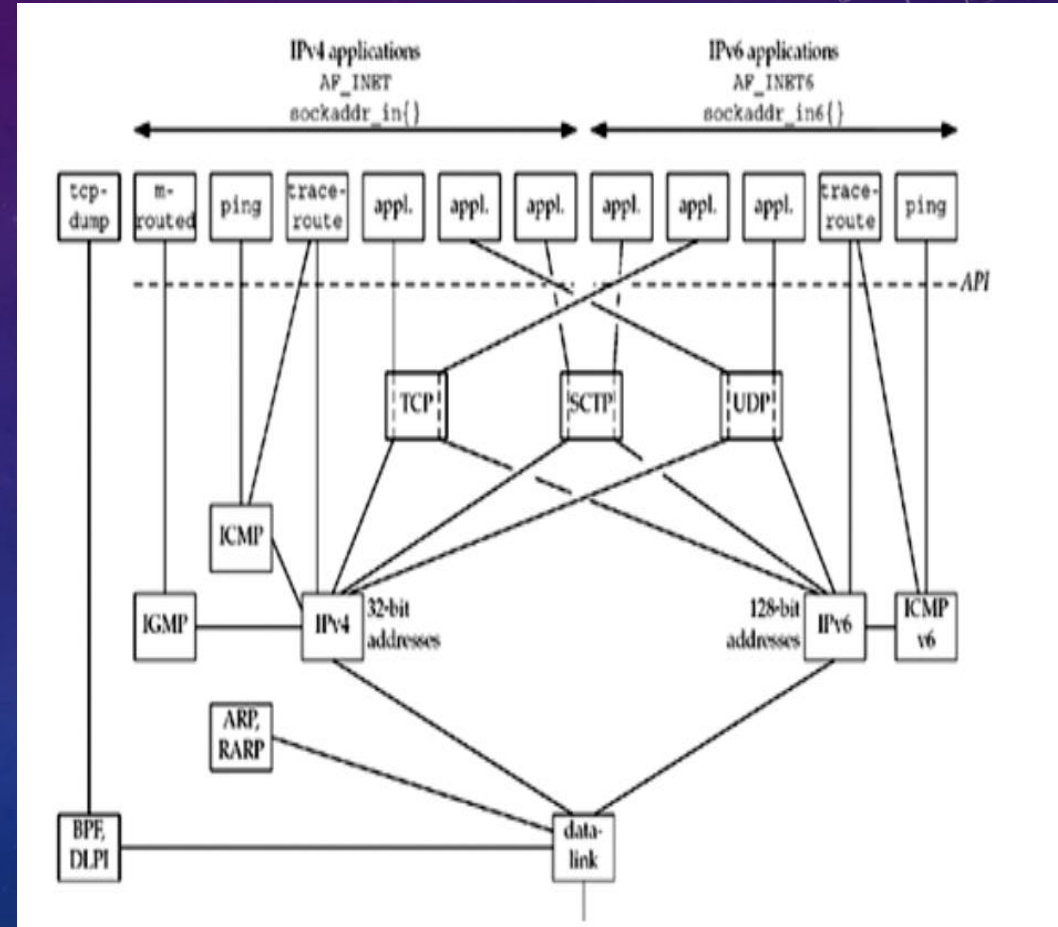
Tin học Nữ Uyên

SOCKET

SOCKET OVERVIEW

- Typical client-server scenario
 - Socket is window to the world
 - Each application open its own socket
 - Server bind its socket to a well-known address
 - Client connect its socket to that address

```
fd = socket(domain, type, protocol);
```



COMMUNICATION DOMAIN

- Communication domain
 - Method of identifying a socket
 - The range of communication (same host or network)
- Type of Domain
 - UNIX : allow communication between application on the same host
 - IPv4: allow communication between applications running on hosts connected via IPv4 network
 - IPv6 allow communication between application running on hosts connected via IPv6

Domain	Communication performed	Communication between applications	Address format	Address structure
AF_UNIX	within kernel	on same host	pathname	<i>sockaddr_un</i>
AF_INET	via IPv4	on hosts connected via an IPv4 network	32-bit IPv4 address + 16-bit port number	<i>sockaddr_in</i>
AF_INET6	via IPv6	on hosts connected via an IPv6 network	128-bit IPv6 address + 16-bit port number	<i>sockaddr_in6</i>

SOCKET TYPES

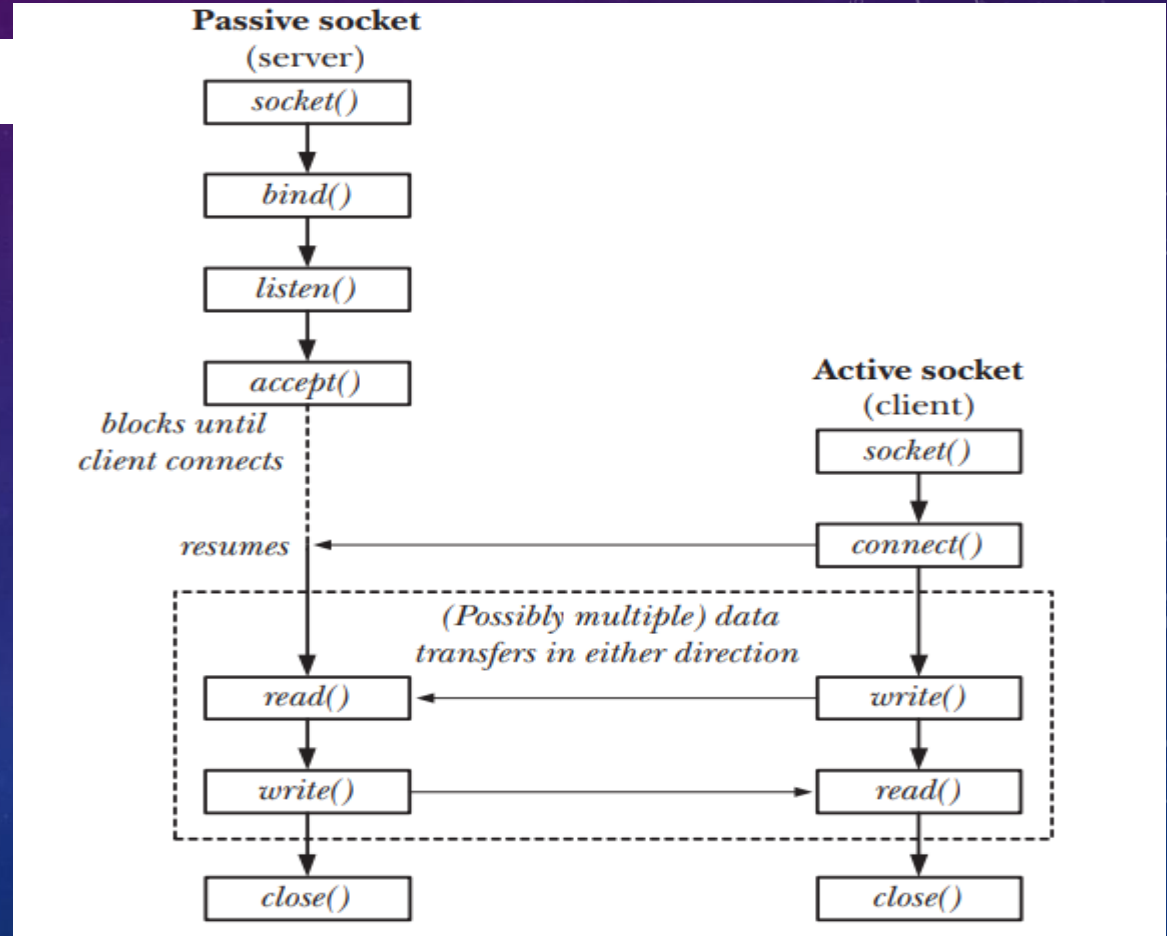
- Two types of socket
 - Stream
 - Reliable – Data at receiver exactly as data at transmitter
 - Bidirectional – Data transmit two directions
 - Byte-stream – No concept of boundary
 - TCP is typical
 - Datagram
 - Exchange in form of message called datagram
 - Message boundaries but data transmission is not reliable , message may be out-of-order, duplicate, not arrived
 - Don't need to establish a connection
 - UDP is typical

Property	Socket type	
	Stream	Datagram
Reliable delivery?	Y	N
Message boundaries preserved?	N	Y
Connection-oriented?	Y	N

SOCKET SYSTEM CALL- STREAM SOCKET FLOW

```
int socket(int domain, int type, int protocol);
```

- Socket
 - Domain AF_UNIX && AF_INET
 - Type SOCK_STREAM && SOCK_DGRAM && SOCK_RAW
 - Protocol IPPROTO_UDP, IPPROTO_TCP, 0
 - Return file descriptor
- Socket operate in connected pair
 - Peer socket – refer to socket at other end of a connection
 - Peer address – denotes the address of peer socket
 - Peer application – denotes application use peer socket



BIND A SOCKET TO AN ADDRESS

```
int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

- Bind a socket to an well-known address
 - Error Address already in use
 - Use `netstat -a` to find which application open it

[illegible]

LISTENING FOR INCOMING CONNECTION

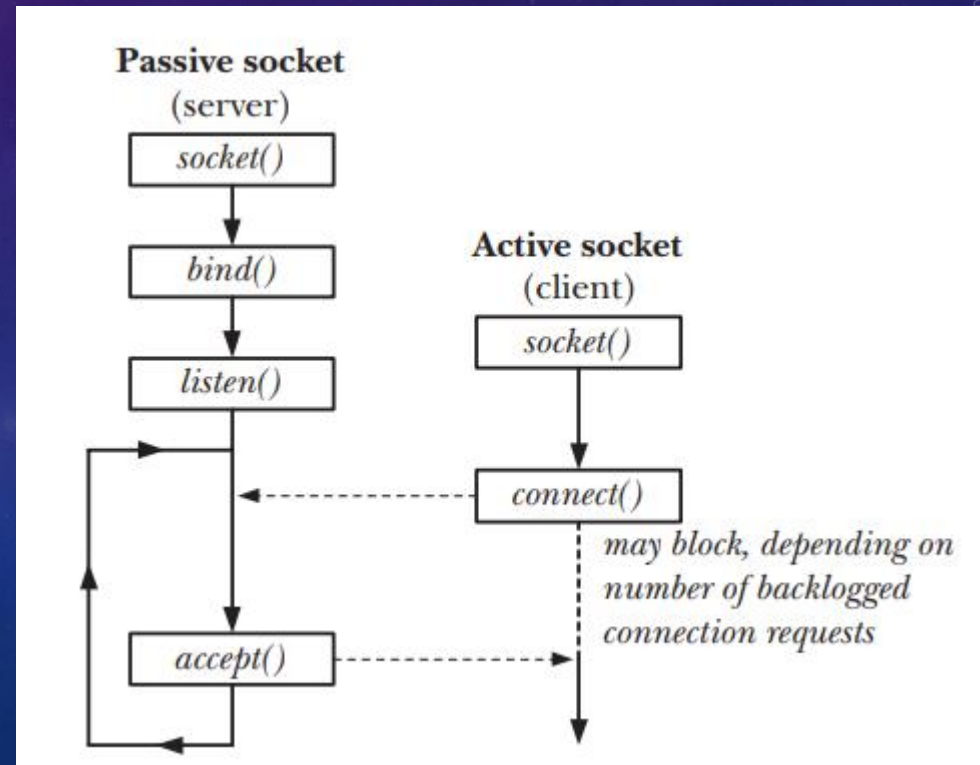
```
int listen(int sockfd, int backlog);
```

- Listen
 - Each connection request coming, kernel record information into pending connection queue → it consume kernel memory
 - `/proc/sys/net/core/somaxconn`
 - Backlog limit number of pending connection
 - `netstat -l` to find all listening port

ACCEPT A CONNECTION

- accept function
 - Accept an incoming connection on listening stream socket refer to by file descriptor
 - If no pending connection exist, accept block until connection request arrive
 - Accept create new socket and keep listening socket to accept new request
 - Addr and addrlen give the address of peer and length of that address
 - Pass NULL to ignore it

```
int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```



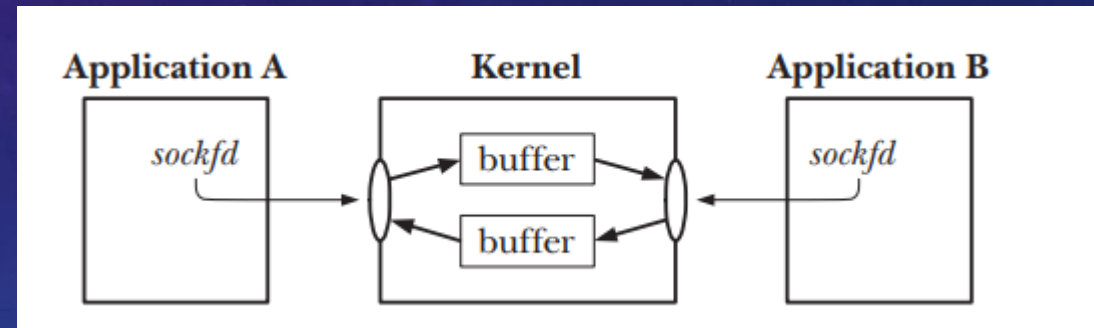
CONNECTING TO PEER SOCKET

- Connect function
 - Connect to a socket refer to by file descriptor whose address specified by `addr` and `addrlen`
 - Establish a connection

```
int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

IO STREAM OPERATION

- Use read/write to request data and send data



TERMINATE A CONNECTION

- Usual way to terminate a connection use close
 - If multiple file descriptor refer to the same socket, the connection is terminated when all of file descriptor are closed
 - If peer crash, with datagram socket, we don't know that event, with stream socket, it have some mechanism to notify that event
 -

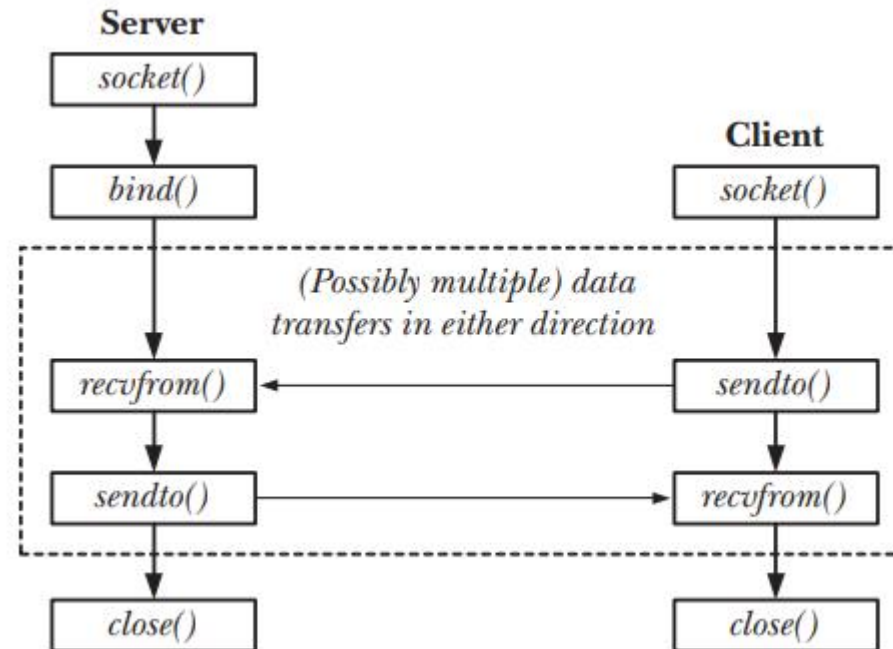
DATAGRAM FLOW

```
ssize_t recvfrom(int sockfd, void *buffer, size_t length, int flags,  
                 struct sockaddr *src_addr, socklen_t *addrlen);
```

Returns number of bytes received, 0 on EOF, or -1 on error

```
ssize_t sendto(int sockfd, const void *buffer, size_t length, int flags,  
               const struct sockaddr *dest_addr, socklen_t addrlen);
```

- sockfd – socket file descriptor
- buffer – start address of message to send
- Length – length of message to send
- Flag – control receive/send as ancillary data, it is usually set to 0
- addr – address of sender in recvfrom and address of receiver in sendto



UNIX DOMAIN SOCKET

- Allow communication between processes on the same host system
- Unix Domain socket addresses
 - Take a form of pathname
 - Use `sprintf` or `strncpy` to avoid buffer overflow of pathname
 - `Ls -l`, we can see type `s` in first column
- Bind a Unix domain socket
 - cast an `sockaddr_un` to `sockaddr`
 - We can't bind an existing pathname, error `EADDRINUSE`
 - Absolute pathname is appreciated
 - A socket can bind only one pathname, a pathname can be bound to only one socket
 - Can't use `open` to open a socket
 - When the socket is no longer require, use `unlink` to remove it

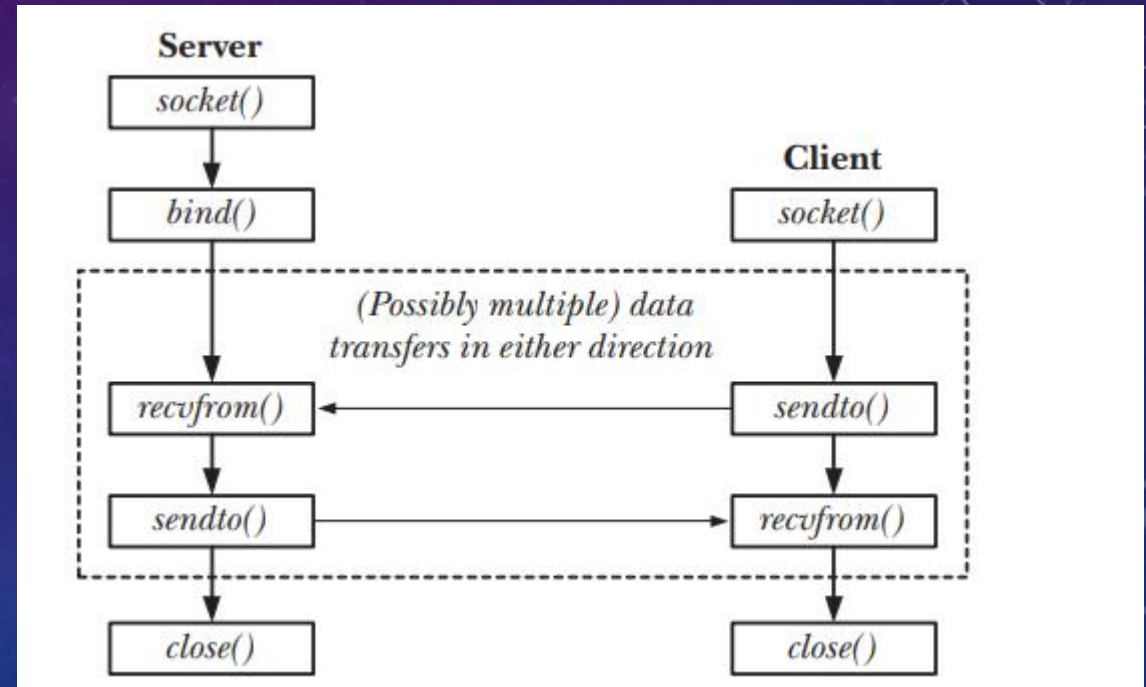
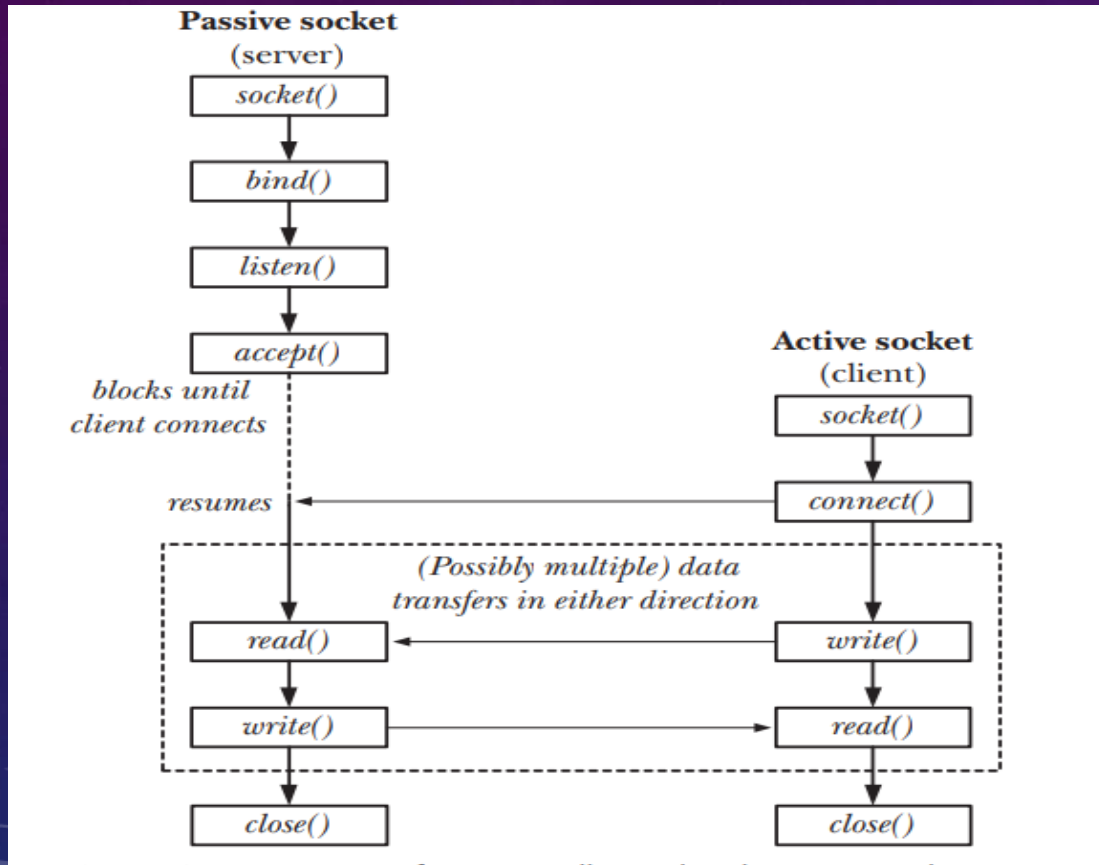
```
struct sockaddr_un {  
    sa_family_t sun_family;    /* Always AF_UNIX */  
    char sun_path[108];        /* Null-terminated socket pathname */  
};
```

```
const char *SOCKNAME = "/tmp/mysock";  
int sfd;  
struct sockaddr_un addr;  
  
sfd = socket(AF_UNIX, SOCK_STREAM, 0);    /* Create socket */  
if (sfd == -1)  
    errExit("socket");  
  
memset(&addr, 0, sizeof(struct sockaddr_un));    /* Clear structure */  
addr.sun_family = AF_UNIX;    /* UNIX domain address */  
strncpy(addr.sun_path, SOCKNAME, sizeof(addr.sun_path) - 1);  
  
if (bind(sfd, (struct sockaddr *) &addr, sizeof(struct sockaddr_un)) == -1)  
    errExit("bind");
```

STREAM SOCKET IN UNIX DOMAIN

- iterative server
 - Create socket
 - Remove/unlink unix pathname
 - Construct an address structure for server socket, bind socket to address and mark socket as listening
 - Loop to handle incoming client request
 - Accept connection, obtain new socket
 - Read data from connected socket and write to standard output
 - Close the socket
 - Server terminate manually
- client
 - create socket
 - Construct an address structure for server socket and connect to socket at address
 - Loop to copies data from standard input to the socket connection until end of file

DATAGRAM SOCKET IN UNIX DOMAIN



MULTIPLEXING IO WITH SELECT

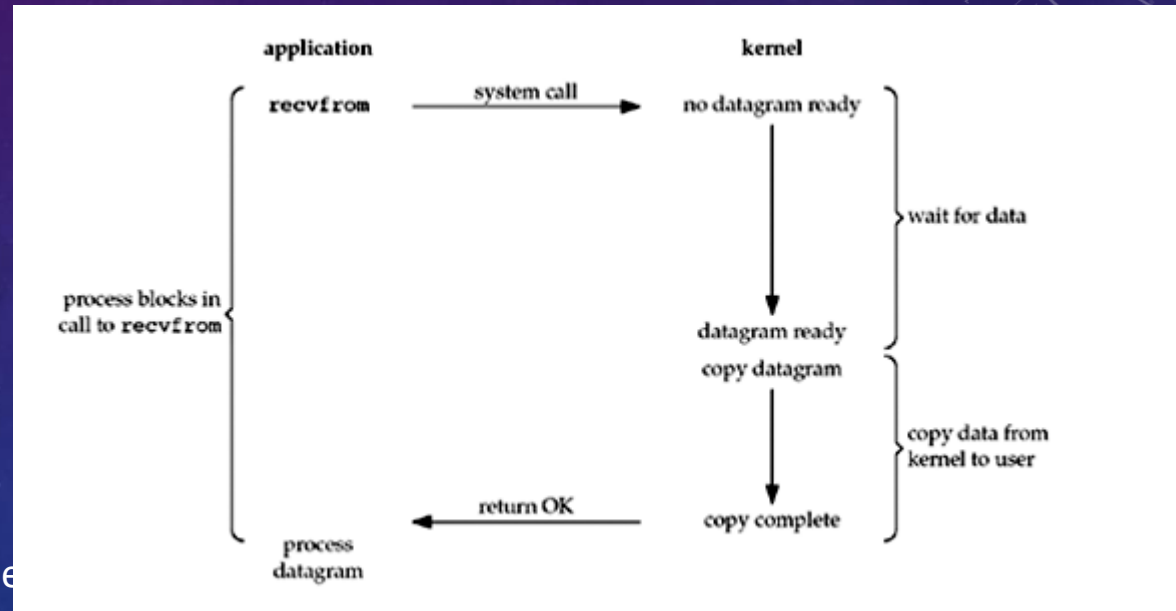
- Server/Client model
 - Server terminate and send goodbye message to client
 - Client busy to waiting command/message from stdio and cannot receive message from server
 - Client have to read from both stdin and socket
 - Client have multiple socket?
- Multiplexing is essential when
 - Client need to handling multiple descriptors(sockets + stdin)
 - TCP handle both a listening socket and connected socket
 - Server handle both TCP(Stream) and UDP(datagram)
 - Server handle multiple service and multiple protocol
 - Not limited to network programming. This technique use every where

I/O MODEL – BLOCKING I/O

- 5 I/O model available
 - Blocking I/O
 - Non-blocking I/O
 - I/O multiplexing
 - Signal driven
 - Asynchronous I/O
- Two phase of input operation
 - Waiting for the data to be ready
 - Copying the data from the kernel to the process

BLOCKING I/O

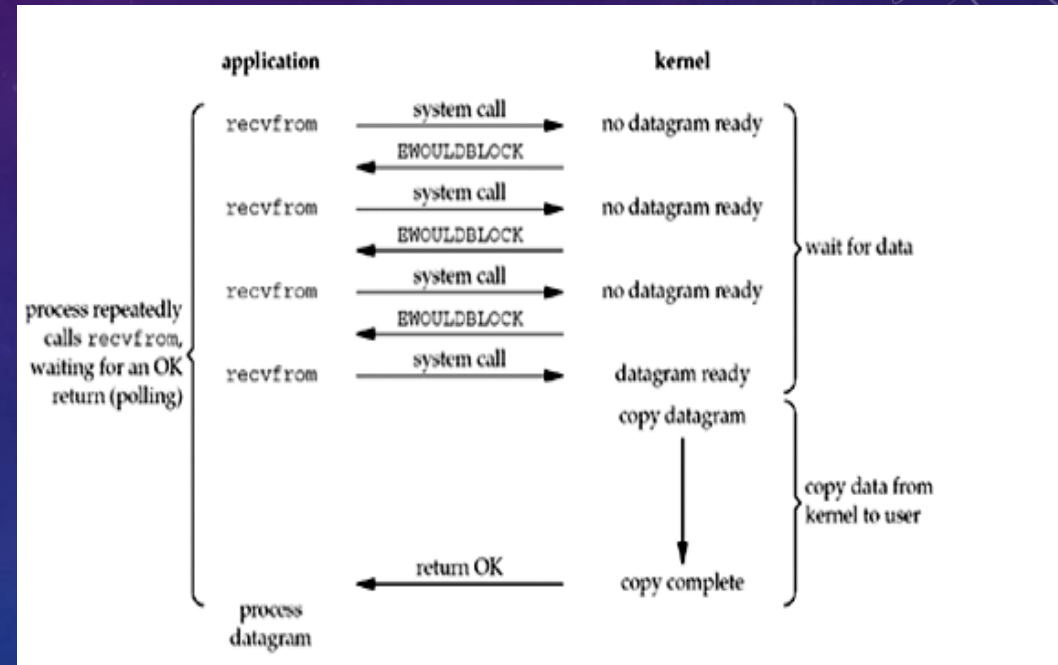
- By default, all sockets are blocking
 - Process waiting a message coming, but message has not arrived, process sleep to wait
 - Calling `recvfrom` cause context-switching and process fire CPU
 - Message is coming, process is wake up to process message
 - `Recvfrom` copy message from kernel buffer to user buffer
 - CPU usage optimization but can't handle multiple events



NON-BLOCKING I/O

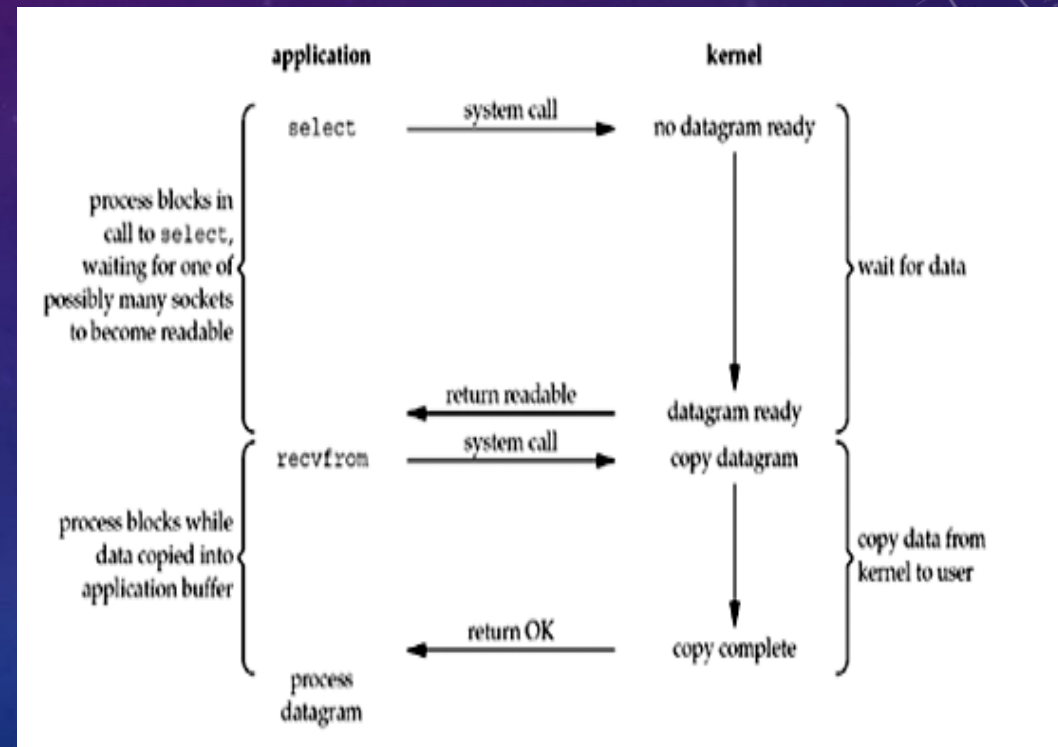
```
int flags = fcntl(fd, F_GETFL, 0);  
fcntl(fd, F_SETFL, flags | O_NONBLOCK);
```

- Socket is set to Non-blocking I/O
 - O_NONBLOCK flags
 - Call recvfrom return immediately with EWOULDBLOCK error code
 - Can handle multiple file descriptor but waste CPU time



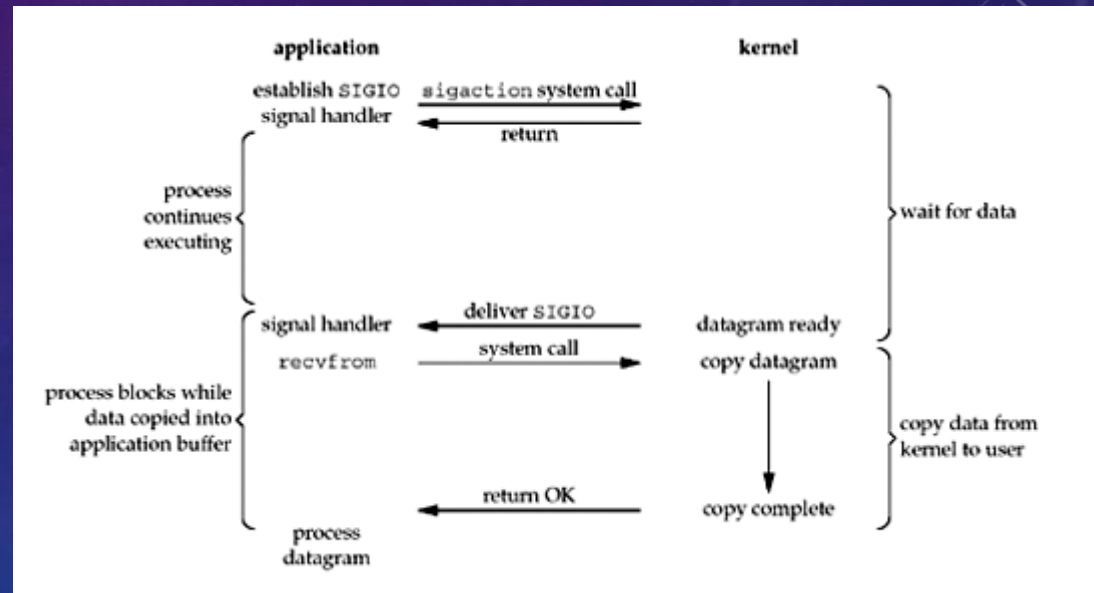
MULTIPLEXING I/O

- I/O multiplexing
 - Call select or polling to wait data from multiple descriptor available
 - Sleep on select instead recvfrom
 - Can handle multiple file descriptor and CPU usage optimization



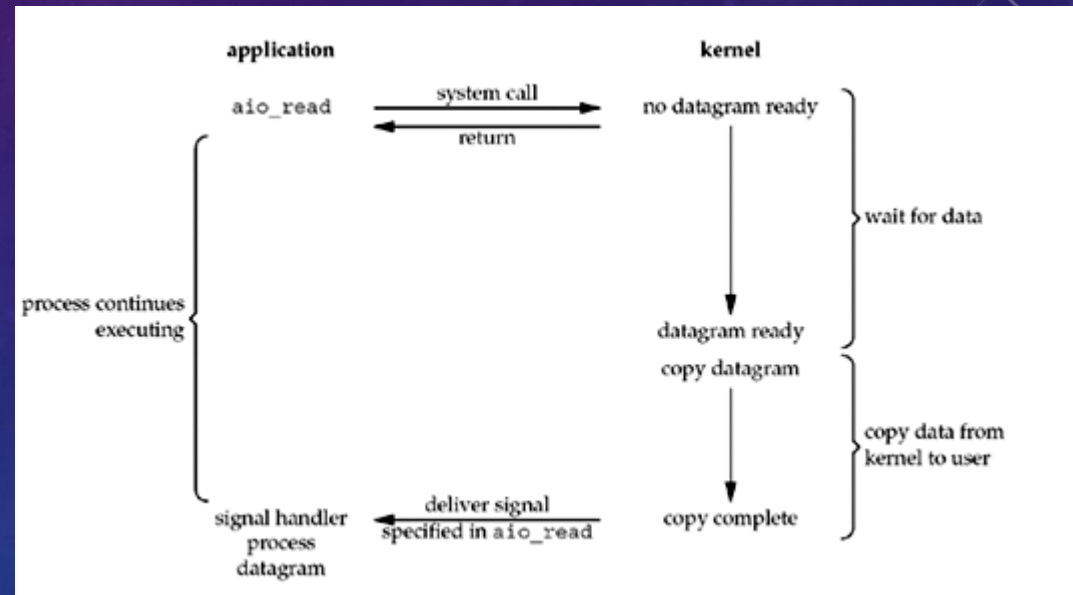
SIGNAL-DRIVEN I/O

- Signal-driven I/O
 - Enable socket for signal-driven I/O
 - Install a signal handler in process
 - When datagram ready to read, a SIGIO signal is generated and deliver to process
 - Can handle multiple file descriptor and CPU usage optimization but it it not popular



ASYNCHRONOUS I/O MODEL

- AIO
 - Call `aio_read` to notify to kernel start operation and notify us when the entire operation is complete
 - The main difference between this model and the signal-driven I/O model is kernel tell process when I/O operation initiated but in AIO, kernel tell us when an I/O operation is completed



SELECT

- Arguments
 - Max of file descriptor tested, its value is file descriptor has maximum value + 1
 - Readset, writeset, exceptionset collect all file descriptor tested (data ready)
 - Timeout - tell the kernel how long we wait the events
- Select blocking until
 - Any events in readset, writeset or exceptionset occur
 - Timer expire
 - File descriptor set/clear/zeroout/test by bitmap operation

```
int select(int nfd, fd_set *readfds, fd_set *writefds,  
          fd_set *exceptfds, struct timeval *timeout);
```

```
void FD_CLR(int fd, fd_set *set);  
int  FD_ISSET(int fd, fd_set *set);  
void FD_SET(int fd, fd_set *set);  
void FD_ZERO(fd_set *set);
```

```
fd_set rset;  
  
FD_ZERO(&rset);           /* initialize the set: all bits off */  
FD_SET(1, &rset);         /* turn on bit for fd 1 */  
FD_SET(4, &rset);         /* turn on bit for fd 4 */  
FD_SET(5, &rset);         /* turn on bit for fd 5 */
```

SEL

- With

-

-

-

- Initia

-

-

- When

```
3 str_cli(FILE *fp, int sockfd)
4 {
5     int      maxfdpl;
6     fd_set   rset;
7     char     sendline[MAXLINE], recvline[MAXLINE];

8     FD_ZERO(&rset);
9     for ( ; ; ) {
10         FD_SET(fileno(fp), &rset);
11         FD_SET(sockfd, &rset);
12         maxfdpl = max(fileno(fp), sockfd) + 1;
13         Select(maxfdpl, &rset, NULL, NULL, NULL);

14         if (FD_ISSET(sockfd, &rset)) { /* socket is readable */
15             if (Readline(sockfd, recvline, MAXLINE) == 0)
16                 err_quit("str_cli: server terminated prematurely");
17             Fputs(recvline, stdout);
18         }

19         if (FD_ISSET(fileno(fp), &rset)) { /* input is readable */
20             if (Fgets(sendline, MAXLINE, fp) == NULL)
21                 return; /* all done */
22             Writen(sockfd, sendline, strlen(sendline));
23         }
24     }
25 }
```

Data to read	•
Read half of the connection closed	•
New connection ready for listening socket	•
Space available for writing	
Write half of the connection closed	
Pending error	•
TCP out-of-band data	